

SQL 4: Joining Tables, Updating and Deleting Data, ACID, and Transactions

Databases and Interfaces

Matthew Pike & Yuan Yao

University of Nottingham Ningbo China (UNNC)

Overview

- In this lecture we will cover:
 - **JOIN**'ing tables in SQL
 - **UPDATE**'ing and **DELETE**'ing data
 - ACID Properties and Transactions

The Database Schema for this Lecture

```
CREATE TABLE Student(  
    sID INTEGER PRIMARY KEY,  
    firstName TEXT NOT NULL,  
    lastName TEXT NOT NULL  
);
```

```
CREATE TABLE Module(  
    mCode CHAR(8) PRIMARY KEY,  
    title TEXT NOT NULL,  
    credits INTEGER NOT NULL  
);
```

```
CREATE TABLE Grade(  
    sID INTEGER NOT NULL,  
    mCode CHAR(8) NOT NULL,  
    grade INTEGER NOT NULL,  
    PRIMARY KEY (sID, mCode),  
    FOREIGN KEY (sID)  
        REFERENCES Student(sID),  
    FOREIGN KEY (mCode)  
        REFERENCES Module(mCode)  
);
```

The Database Content for this Lecture

sID	firstName	lastName
1	John	Smith
2	Jane	Doe
3	Mary	Jones
4	Joe	Bloggs

Table 1: Student Table

mCode	title	credits
COMP1036	Fundamentals	20
COMP1048	Databases	10
COMP1038	Programming	20

Table 2: Module Table

sID	mCode	grade
1	COMP1036	35
1	COMP1048	50
2	COMP1048	65
2	COMP1038	70
3	COMP1036	35
3	COMP1038	65
6	COMP1038	55
6	COMP1099	68

Table 3: Grade Table

Joining Tables in SQL

- JOINS can be used to combine rows from two or more tables, based on a related column between them.
- This is a very common operation in relational databases, as it allows you to link information between different tables.
- There are numerous types of JOINS, but the most common are:
 - CROSS JOIN
 - INNER JOIN
 - LEFT & RIGHT JOIN
 - NATURAL JOIN
 - FULL OUTER JOIN

CROSS JOIN

- The **CROSS JOIN** returns the Cartesian product of the sets of rows from the two tables.
 - This means that every row from the first table is combined with every row from the second table.
 - This also means that the resulting table will contain rows of data that are not related (nonsensical data).
- The syntax for a **CROSS JOIN** is:
 - **SELECT * FROM table1 CROSS JOIN table2;**
 - Which is equivalent to:
 - **SELECT * FROM table1, table2;**
- **CROSS JOIN** is rarely used in practice, as it can result in a very large number of rows.
 - We can constrain the number of rows returned by using a **WHERE** clause.

Example: CROSS JOIN

```
SELECT * FROM Student CROSS JOIN Module LIMIT 8;
```

sID	firstName	lastName	mCode	title	credits
1	John	Smith	COMP1036	Fundamentals	20
1	John	Smith	COMP1048	Databases	10
1	John	Smith	COMP1038	Programming	20
2	Jane	Doe	COMP1036	Fundamentals	20
2	Jane	Doe	COMP1048	Databases	10
2	Jane	Doe	COMP1038	Programming	20
3	Mary	Jones	COMP1036	Fundamentals	20
3	Mary	Jones	COMP1048	Databases	10

Table 4: The first 8 results of the CROSS JOIN of Student and Module

SELECT from Multiple Tables

- **SELECT** can be used with multiple tables, with table names separated by commas in the **FROM** clause.
 - `SELECT * FROM Student, Module;`
 - This is equivalent to a **CROSS JOIN** of the two tables.
- We can limit the columns returned by **SELECT** by specifying the table name before the column name.
 - `SELECT Student.sID, Module.mCode FROM Student, Module;`

Example: SELECT from Multiple Tables

```
SELECT
Student.sID,
Module.mCode,
grade --Not ambiguous
FROM
    Student, Grade, Module
WHERE
    Student.sID = Grade.sID
    AND
    Module.mCode = Grade.mCode;
```

sID	mCode	grade
1	COMP1036	35
1	COMP1048	50
2	COMP1048	65
2	COMP1038	70
3	COMP1036	35
3	COMP1038	65

Table 5: The SELECT from Multiple Tables

CROSS JOIN不可以用ON Clause

- The **INNER JOIN** returns only rows where the join condition is met.
- The join condition is specified in the **ON clause**.
 - `SELECT * FROM table1 INNER JOIN table2 ON table1.column1 = table2.column2;`

Example: INNER JOIN

```
SELECT
    Student.lastName,
    Grade.grade
FROM
    Student INNER JOIN Grade
ON
    Student.sID = Grade.sID;
```

lastName	grade
Smith	35
Smith	50
Doe	65
Doe	70
Jones	35
Jones	65

Table 6: The INNER JOIN of Student and Grade

LEFT JOIN

- The **LEFT JOIN** returns all rows from the left table, and the matched rows from the right table.
- Any rows from the right table that do not have a match in the left table are returned with **NULL** values.
- Left joins are often used when you want to see all the rows from one table, even if there is no match in the other table.
- The syntax for a **LEFT JOIN** is:
 - `SELECT * FROM table1 LEFT JOIN table2 ON condition;`

LEFT JOIN Example

```
SELECT
    Student.sID,
    Student.lastName AS "Last",
    Module.mCode AS "Module",
    Grade.grade as 'Grade'
FROM
    Student LEFT JOIN Grade
ON
    Student.sID = Grade.sID
LEFT JOIN Module
ON
    Grade.mCode = Module.mCode;
```

sID	Last	Module	Grade
1	Smith	COMP1036	35
1	Smith	COMP1048	50
2	Doe	COMP1038	70
2	Doe	COMP1048	65
3	Jones	COMP1036	35
3	Jones	COMP1038	65
4	Bloggs	NA	NA

Table 7: The LEFT JOIN of Student and Grade. Note the final row.

! RIGHT JOIN support in SQLite

Support for **RIGHT JOIN** is only available in SQLite version 3.39.0 and above.

- The **RIGHT JOIN** returns all rows from the right table, and the matched rows from the left table.
- Any rows from the left table that do not have a match in the right table are returned with **NULL** values.
- The syntax for a **RIGHT JOIN** is:
 - `SELECT * FROM table1 RIGHT JOIN table2 ON condition;`

RIGHT JOIN Example

```
SELECT
    Student.sID,
    Student.lastName AS "Last",
    Module.mCode AS "Module",
    Grade.grade as 'Grade'
FROM
    Student RIGHT JOIN Grade
ON
    Student.sID = Grade.sID
LEFT JOIN Module
ON
    Grade.mCode = Module.mCode;
```

sID	Last	Module	Grade
1	Smith	COMP1036	35
1	Smith	COMP1048	50
2	Doe	COMP1038	70
2	Doe	COMP1048	65
3	Jones	COMP1036	35
3	Jones	COMP1038	65
NA	NA	COMP1038	55
NA	NA	NA	68

Table 8: The RIGHT JOIN of Student and Grade.

- The **NATURAL JOIN** returns all rows where the join condition is met.
- The syntax for a **NATURAL JOIN** is:
 - `SELECT * FROM table1 NATURAL JOIN table2;`
- The **NATURAL JOIN** can only be used if the columns to be joined have the same name in both tables.

Example: NATURAL JOIN

```
SELECT
    Student.sID,
    Student.lastName AS "Last",
    Module.mCode AS "Module",
    Grade.grade as 'Grade'
FROM
    Student
    NATURAL JOIN -- sID
    Grade
    NATURAL JOIN -- mCode
    Module;
```

sID	Last	Module	Grade
1	Smith	COMP1036	35
1	Smith	COMP1048	50
2	Doe	COMP1048	65
2	Doe	COMP1038	70
3	Jones	COMP1036	35
3	Jones	COMP1038	65

Table 9: The NATURAL JOIN of Student, Grade, and Module

FULL OUTER JOIN

! FULL OUTER JOIN support in SQLite

Support for FULL OUTER JOIN is only available in SQLite version 3.39.0 and above.

LEFT和RIGHT取交集，NULL都显示出来

- The FULL OUTER JOIN returns all rows from both tables, where the join condition is met.
- Any rows from the left table that do not have a match in the right table are returned with NULL values.
- Any rows from the right table that do not have a match in the left table are returned with NULL values.
- The syntax for a FULL OUTER JOIN is:
 - `SELECT * FROM table1 FULL OUTER JOIN table2 ON condition;`

Example: FULL OUTER JOIN

```
SELECT
    Student.sID,
    Student.lastName AS "Last",
    Module.mCode AS "Module",
    Grade.grade as 'Grade'
FROM
    Student FULL OUTER JOIN Grade
ON
    Student.sID = Grade.sID
FULL OUTER JOIN Module
ON
    Grade.mCode = Module.mCode;
```

sID	Last	Module	Grade
1	Smith	COMP1036	35
1	Smith	COMP1048	50
2	Doe	COMP1038	70
2	Doe	COMP1048	65
3	Jones	COMP1036	35
3	Jones	COMP1038	65
4	Bloggs	NA	NA
NA	NA	COMP1038	55
NA	NA	NA	68

Table 10: The FULL OUTER JOIN of Student, Grade, and Module

Updating Data in SQL

UPDATE Statement

- The **UPDATE** statement is used to modify the existing records in a table.
- The syntax for the **UPDATE** statement is:
 - `UPDATE table_name SET column1 = value1, column2 = value2, ... WHERE condition;`
- The **WHERE** clause is optional.
 - If it is omitted, all records in the table will be updated.
- Within the **SET** clause, you can specify multiple columns and values.
- The **UPDATE** statement can reference column values from other columns in the same row.
 - For example, `UPDATE table SET column1 = column1 + 1;` will increment the value of `column1` by 1.

Example: UPDATE Statement

```
UPDATE Student
SET
    firstName = 'Johnathan',
    lastName = 'Creek'
WHERE sID = 1;
```

```
SELECT * FROM Student;
```

sID	firstName	lastName
1	Johnathan	Creek
2	Jane	Doe
3	Mary	Jones
4	Joe	Bloggs

Table 11: The Student table after UPDATE

Deleting Data with SQL

DELETE Statement

A Note on Refential Integrity

Remember, tables with foreign keys have a **refential integrity** constraint, which means that the **DELETE statement may fail if the foreign key is referenced in another table, unless the CASCADE option is used.** However, by default, SQLite does not enforce referential integrity. To enable it, we need to use the **PRAGMA** statement: **PRAGMA foreign_keys = ON;**.

- The **DELETE** statement is used to delete existing records in a table.
- The syntax for the **DELETE** statement is:
 - **DELETE FROM table_name WHERE condition;**
- The **WHERE** clause is optional.
 - If it is omitted, all records in the table will be deleted.
- The **DELETE** statement returns the number of rows that were deleted.

Example: DELETE Statement

```
DELETE FROM Student WHERE sID = 4;
```

```
DELETE FROM Grade WHERE sID = 6;
```

```
SELECT *  
FROM Student  
WHERE sID >= 3;
```

sID	firstName	lastName
3	Mary	Jones

Table 12: The Student table after DELETE

```
SELECT *  
FROM Grade  
WHERE sID >= 4;
```

sID	mCode	grade

Table 13: The Grade table after DELETE

Transactions

- ACID is an acronym for the four properties of transactions:
 - **Atomicity:** All or nothing. Either all of the operations in a transaction are completed, or none of them are.
 - **Consistency:** The database is always in a valid state.
 - **Isolation:** Transactions are isolated from each other.
 - **Durability:** Once a transaction has been committed, it will remain so, even in the event of a system failure.
- SQLite is ACID compliant and supports transactions.
- SQLite guarantees that all transactions are ACID compliant even if the transaction is interrupted by a power failure or system crash.

- A transaction is a sequence of SQL statements that are treated as a single unit.
 - Either all of the statements are executed, or none of them are.
- Transactions are used to ensure that the database is in a consistent state after the transaction is completed.
 - For example, if a transaction updates two tables, and one of the updates fails, the database should be left in the same state as before the transaction was started.

- The syntax for a transaction is:
 - **BEGIN TRANSACTION;**
 - `-- SQL statements`
 - **COMMIT;**
- The **BEGIN TRANSACTION** statement starts a transaction.
- The **COMMIT** statement commits the transaction, which means that the changes are saved to the database.
- If any of the SQL statements in the transaction fail, the **ROLLBACK** statement can be used to undo the changes.
 - **ROLLBACK;**

Example: Transactions

```
BEGIN TRANSACTION;  
    INSERT INTO Student VALUES (4, 'John', 'Doe');  
    INSERT INTO Student VALUES (5, 'Jane', 'Smith');  
    INSERT INTO Student VALUES (6, 'John', 'Smith');  
  
-- Commit the changes to the database:  
COMMIT;  
  
-- If you do not want to save the changes:  
ROLLBACK;
```


References

Online Tutorials

These are clickable links to the online tutorials:

- [Join Operators](#)
- [Update](#)
- [Delete](#)
- [Transactions](#)
- [A Visual Explanation of SQL Joins](#)

Textbooks and Documentation

- [Chapter 5 and 22 of the Databases textbook.](#)
- [SQLite Transactions](#)
- [SQLite Joins](#)

Mohan, Chandrasekaran, Don Haderle, Bruce Lindsay, Hamid Pirahesh, and Peter Schwarz. 1992. "ARIES: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using Write-Ahead Logging." *ACM Transactions on Database Systems (TODS)* 17 (1): 94–162.