

# COMP2005

---

Introduction to Image Compression

# Image Compression

- Individual images can be large
- Images are easy to acquire, collections increase rapidly
- In some applications, images are gathered automatically
- Luckily, image data is redundant in several ways
  - Coding redundancy
  - Spatial redundancy
  - Psychovisual redundancy



# Coding Redundancy

- The grey level histogram of an image gives the probability (frequency) of occurrence of grey level  $r_k$

$$p(r_k) = \frac{n_k}{n} \quad k = 0, 2, \dots, L-1$$

- If the number of bits used to represent each value of  $r_k$  is  $l(r_k)$ , the **average number of bits required to represent a pixel is**

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k)$$

- To code an  $M \times N$  image requires  $MNL_{avg}$  bits

# Coding Redundancy

- If an m-bit natural binary code is used to represent grey level then
  - all pixels take the same amount of space,
  - $P(r_k)$  values sum to 1, so

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k) = \sum_{k=0}^{L-1} m p(r_k) = m$$

- and an image occupies MNm bits
- But some pixel values are more common than others.....

# Variable Length Encoding

- Assigning fewer bits to the more probable grey levels than to less probable ones can achieve data compression, e.g:

$r_k$	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
$r_k$ for $k \neq 87, 128, 186, 255$	0	—	8	—	0


- Build a codebook, replace 'true' pixel values with code
- Lossless: the process can be reversed by inverting the codebook

# Spatial Redundancy

- Sometimes called *Interpixel Redundancy*
- Neighbouring pixels often have similar values
- Compression based on spatial redundancy involves some element of pixel grouping, or transformation
- Simplest is Run-Length Encoding
  - maps the pixels along each scan line into a sequence of pairs  $(g_1, r_1), (g_2, r_2), \dots,$
  - where  $g_i$  is the  $i$ th grey level,  $r_i$  is the run length of  $i$ th run

# A Binary Example

Row 1: (0, 16)  
 Row 2: (0, 16)  
 Row 3: (0, 7) (1, 2) (0, 7) ← encode  
 Row 4: (0, 4), (1, 8) (0, 4) ←  
 Row 5: (0, 3) (1, 2) (0, 6) (1, 3) (0, 2)  
 Row 6: (0, 2) (1, 2) (0, 8) (1, 2) (0, 2)  
 Row 7: (0, 2) (1, 1) (0, 10) (1, 1) (0, 2)  
 Row 8: (1, 3) (0, 10) (1, 3)  
 Row 9: (1, 3) (0, 10) (1, 3)  
 Row 10: (0, 2) (1, 1) (0, 10) (1, 1) (0, 2)  
 Row 11: (0, 2) (1, 2) (0, 8) (1, 2) (0, 2)  
 Row 12: (0, 3) (1, 2) (0, 6) (1, 3) (0, 2)  
 Row 13: (0, 4) (1, 8) (0, 4)  
 Row 14: (0, 7) (1, 2) (0, 7)  
 Row 15: (0, 16)  
 Row 16: (0, 16) → decode



# Psychovisual Redundancy

- Some grey level and colour differences are imperceptible; goal is to compress without noticeable change to the image

256 gray levels



16 gray levels



16 gray levels



A simple method:  
add a small  
random number to  
each pixel before  
quantization



# Evaluating Compression

- Fidelity Criteria: success is judged by comparing original and compressed versions
- Some measures are objective, e.g. root mean square error ( $e_{rms}$ ) and signal to noise ratio (SNR)
- Let  $f(x,y)$  be the input image,  $f'(x,y)$  be reconstructed input image from compressed bit stream, then

$$e_{rms} = \left( \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y) - f(x,y))^2 \right)^{1/2}$$

$$SNR = \frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y))^2}{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y) - f(x,y))^2}$$

# Fidelity Criteria



$$e_{rms} = 6.93$$

$$SNR_{rm} = 10.25$$

$$e_{rms} = 6.78$$

$$SNR_{rm} = 10.39$$

# Fidelity Criteria

- $e_{\text{rms}}$  and SNR are convenient objective measures
- Most decompressed images are viewed by human beings
- Subjective evaluation of compressed image quality by human observers are often more appropriate


Value	Rating	Description
1	Excellent	An image of extremely high quality, as good as you could desire.
2	Fine	An image of high quality, providing enjoyable viewing. Interference is not objectionable.
3	Passable	An image of acceptable quality. Interference is not objectionable.
4	Marginal	An image of poor quality; you wish you could improve it. Interference is somewhat objectionable.
5	Inferior	A very poor image, but you could watch it. Objectionable interference is definitely present.
6	Unusable	An image so bad that you could not watch it.

Rating scale of the Television Allocations Study Organization. (Freundtall and Behrend.)

# Image Compression Systems

- Mapper (spatial redundancy)
  - transforms input data in a way that facilitates reduction of interpixel redundancies
  - reversible
- Quantiser (psychvisual redundancy)
  - transforms input data in a way that facilitates reduction of psychovisual redundancies
  - **not reversible**
- Symbol coder (coding redundancy)
  - assigns the shortest code to the most frequently occurring output values
  - reversible

# Image Compression Systems



Functional block diagram of a general image compression system

---

## Image Compression: Huffman Coding

# Exploiting Coding Redundancy

- These methods are derived from information theory:  
**try to maintain a high level of information in compressed images**
- Not limited to images, are applicable to any digital information.
  - speak of ***symbols*** instead of ***pixel values*** and ***sources*** instead of ***images***
  - exploit nonuniform probabilities of symbols (nonuniform histograms) and use a variable-length code.
- Evaluation requires a measure of the information content of a source: *Entropy*

# Entropy

- The idea: associate information with probability
- A random event  $E$  with probability  $P(E)$  contains:

$$I(E) = \log\left(\frac{1}{P(E)}\right) = -\log(P(E))$$

units of information

- Suppose that grey level values are generated by a random variable, then  $r_k$  contains:

$$I(r_k) = -\log(P(r_k))$$


Note:  $I(E)=0$  when  $P(E)=1$

units of information





# Entropy

- Entropy is the average information content of an image, a measure of histogram dispersion



$$H = -\sum_{k=0}^{L-1} p(r_k) \log_2 p(r_k)$$

*entropy=7.4635*



*entropy=0*

- Can't compress to less than H bits/pixel without losing information

# Huffman Coding

- Compute probabilities of each symbol by histogramming the source
- Process probabilities to pre-compute codebook:  
code(i)
  - codebook is static (fixed)
- Encode source symbol-by-symbol:  
symbol(i)  $\rightarrow$  code(i)
- Transmit coded signal and codebook
- The need to pre-process (histogram) the source before encoding begins is a disadvantage

# Huffman Coding

- Builds a binary tree in which symbols to be coded are nodes
- The algorithm:
  - Create a list of nodes, one per symbol, sorted in order of symbol frequency (or probability)
  - REPEAT (until only one node left)
    - Pick the two nodes with the lowest frequencies/probabilities and create a parent of them
    - **Randomly** assign the codes 0,1 to the two new branches of the tree and delete the children from the list
    - Assign the sum of the children's probabilities to their parent and insert it in the list
- Path from root to node gives code for corresponding symbol

# Huffman Coding


Original source		Source reduction			
Symbol	Probability	1	2	3	4
$a_2$	0.4	0.4	0.4	0.4	0.6 0.4
$a_6$	0.3	0.3	0.3	0.3	
$a_1$	0.1	0.1	0.2	0.3	
$a_4$	0.1	0.1	0.1		
$a_3$	0.06	0.1			
$a_5$	0.04				

- Create parent node of  $a_3$  and  $a_5$

- Create parent of new node and  $a_4$

- New parent is higher up list (0.2) than  $a_1$  (0.1)

# Huffman Coding



Original source			Source reduction			
Symbol	Probability	Code	1	2	3	4
a <sub>2</sub>	0.4	1	0.4	1	0.4	1
a <sub>6</sub>	0.3	00	0.3	00	0.3	00
a <sub>1</sub>	0.1	011	0.1	011	0.2	010
a <sub>4</sub>	0.1	0100	0.1	0100	0.1	011
a <sub>3</sub>	0.06	01010	0.1	0101		
a <sub>5</sub>	0.04	01011				

# Huffman Coding

- The algorithm systematically places nodes representing high probability symbols further up the tree: their paths (and so codes) are shorter
- No code is a prefix to any other – don't need to mark boundaries between codes
  - e.g. 01101010 must be  $a_1a_3$
- In this example
  - Average length of the code is 2.2.bits/symbol
  - The entropy of the source is 2.14 bits/symbol
- In general
  - Break image into small (e.g. 8 x 8) blocks
  - Each block is a symbol to be encoded

# A Huffman Code Example

From a past exam paper:

*An image has the following normalized histogram. Derive a Huffman code for each pixel value, showing how you obtained your code*

10 mins

Pixel Value	Normalised Frequency
0	0.1
1	0.1
2	0.15
3	0.35
4	0.2
5	0
6	0.05
7	0.05

# A Huffman Code Example

Pixel Value	Normalised Frequency
0	0.1
1	0.1
2	0.15
3	0.35
4	0.2
5	0
6	0.05
7	0.05

Sort




Pixel Value	Normalised Frequency
3	0.35
4	0.2
2	0.15
1	0.1
0	0.1
7	0.05
6	0.05
5	0



# A Huffman Code Example: Table View

Pixel Value	Normalised Frequency							
3	0.35	0.35	0.35	0.35	0.35	0.4	0.6	1.0
4	0.2	0.2	0.2	0.2	0.25	0.35	0.4	
2	0.15	0.15	0.15	0.2	0.2	0.25		
1	0.1	0.1	0.1	0.15	0.2			
0	0.1	0.1	0.1	0.1				
7	0.05	0.05	0.1					
6	0.05	0.05						
5	0							

# A Huffman Code Example: Tree View



# A Huffman Code Example: The Code

		Code length	Normalised Freq,
3: 00	0: 110	3	0.1
4: 10	1: 011	3	0.1
2: 010	2: 010	3	0.15
1: 011	3: 00	2	0.35
0: 110	4: 10	2	0.2
7: 1110	5: 11111	5	0
6: 11110	6: 11110	5	0.05
5: 11111	7: 1110	4	0.05

Mean bits/pixel  $L_{avg} = 0.3 + 0.3 + 0.45 + 0.70 + 0.4 + 0 + 0.25 + 0.2 = 2.6$

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k)$$

Without compression  $L_{avg} = 3$

Compression ratio =  $2.6/3 = 0.86$


---

## Image Compression: GIF


# Using Psychovisual Redundancy

- Quantisation is widely used
  - Represent areas of grey level/colour space with fewer bits
  - Lossy: cannot be inverted
  - Find the best tradeoff between  
maximal compression  $\leftrightarrow$  minimal distortion
- Scalar Quantisation (i.e. quantising scalar values)

Uniform scalar  
quantization:



Non-uniform scalar  
quantization:



# Vector Quantisation

- Palettised images (gif)
  - Map vector values (R,G,B) onto scalar values
  - Multiple vectors map to each scalar



True colour R,G,B  
8 bits each  
1677216 possible colours




Vector  
quantisation




gif  
8 bits per pixel  
256 possible colours

# Paletised Images




# Building a Palette






# Building a Palette



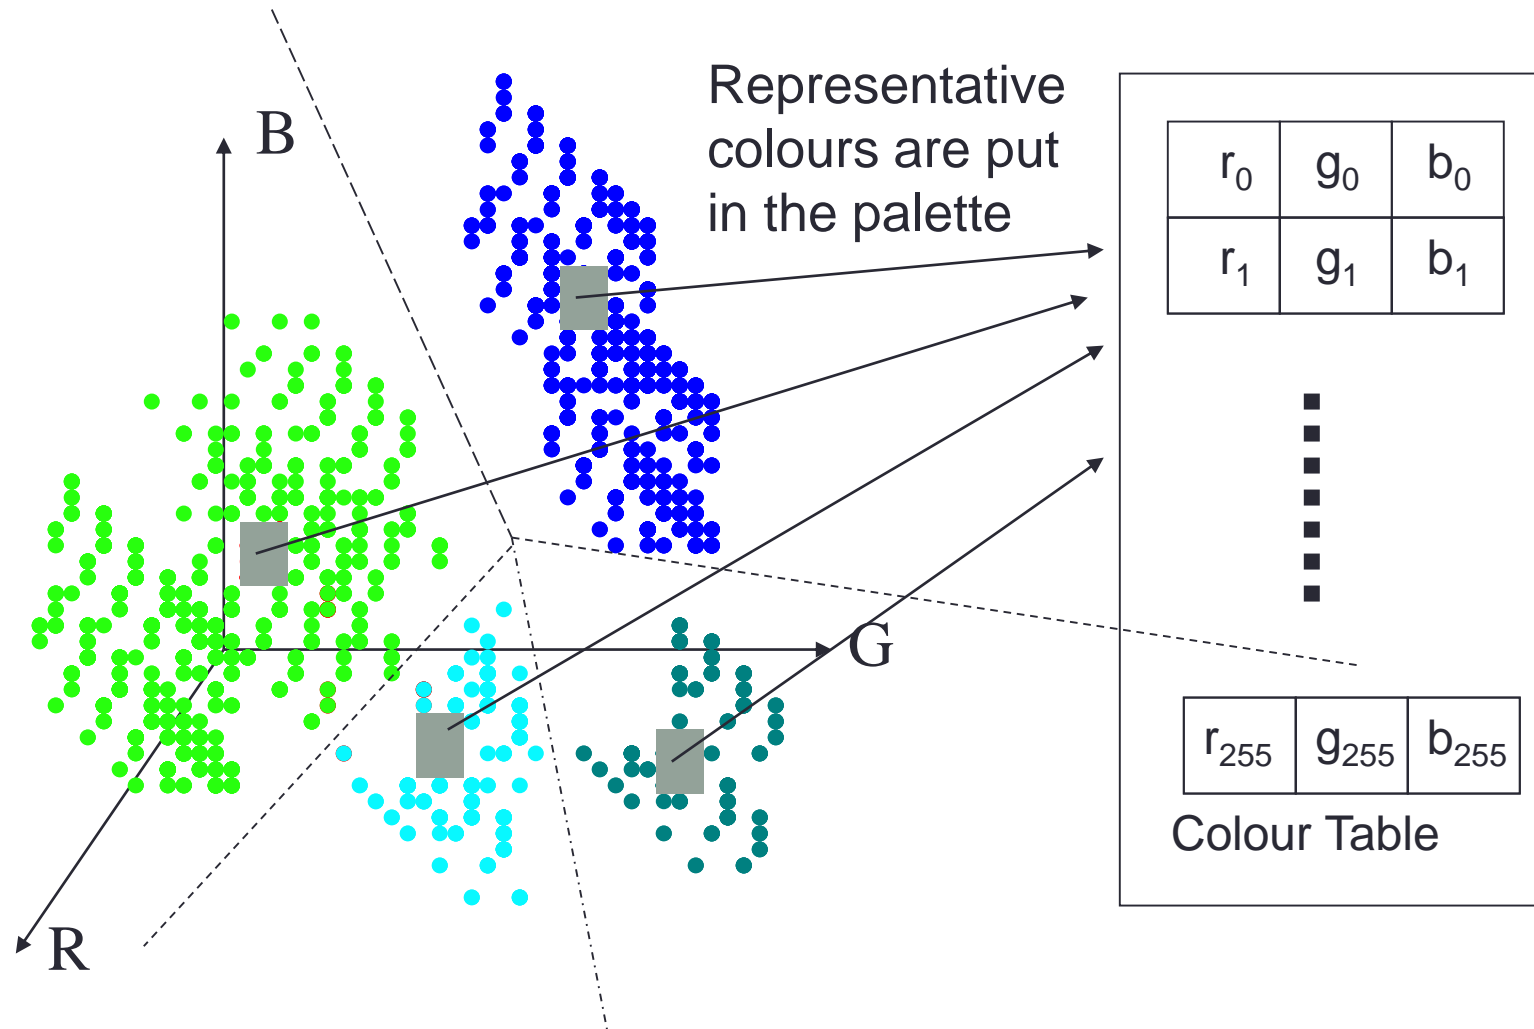
# Building a Palette



Group pixels that are close to each other, and replace them by one single colour

Close to each other means of a similar colour

# Building a Palette



# Building a Palette

- Many clustering algorithms exist
  - supervised
  - unsupervised
- We know how many clusters we need: one per palette entry
- We need clusters that are spread across the colour space
- A supervised method....
- K-Means Clustering
  - Start with estimates of the mean of each cluster
$$\mu_1, \mu_2, \dots, \mu_k$$
  - Assign each point,  $p$ , to the cluster where
$$|p - \mu_i|$$
is smallest
  - Recompute the means
  - Repeat until no changes are made to the clusters



# COMP2005

---

Image Compression: JPEG



# Exploiting Spatial Redundancy


- Run-length encoding needs adjacent pixels to be **equal**
- Pixels are more often **highly correlated** (dependent)
  - Not equal, but can predict the image pixels to be coded from those already coded
- Each pixel value (except at the boundaries) is predicted based on its neighbors (e.g., linear combination) to get a **predicted** image.
- The difference between the original and predicted images yields a **differential** or **residual** image with a reduced set of values.
- The differential image is encoded using Huffman coding, or similar.

# Differential Pulse-code Modulation

- Code the difference between adjacent pixels

Original pixels: 82, 83, 86, 88, 56, 55, 56, 60, 58, 55, 50, .....  $\longrightarrow$  DPCM: 82, 1, 3, 2, -32, -1, 1, 4, -2, -3, -5, .....

- Prediction is that the next pixel value = current one
- Need the first value to provide a point of reference
- Invertible (lossless) and lower entropy





# Predictive Coding

- Higher order pattern prediction
- Use both 1D and 2D patterns (to predict shaded pixel)

1D Causal:





2D Causal:



1D Non-causal:



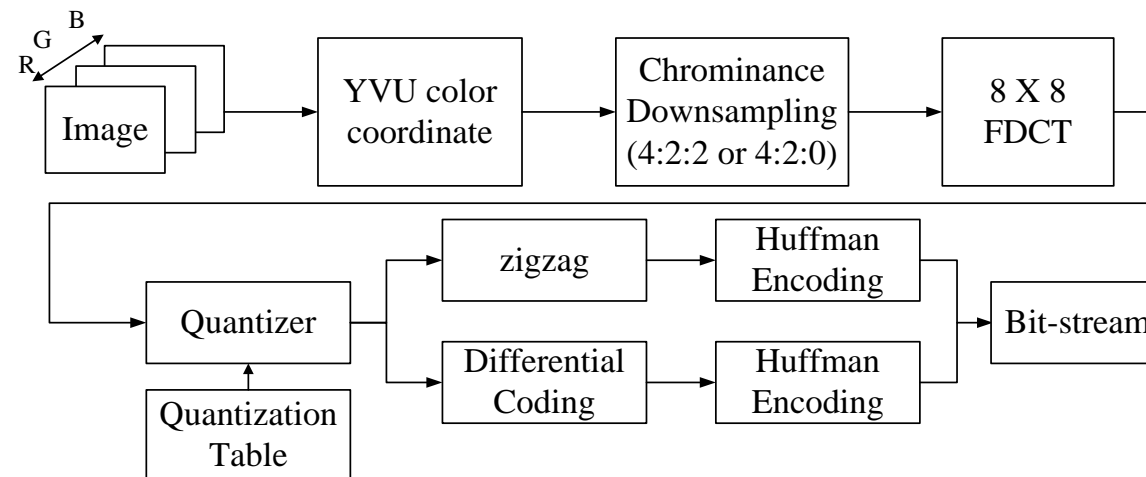
2D Non-Causal:





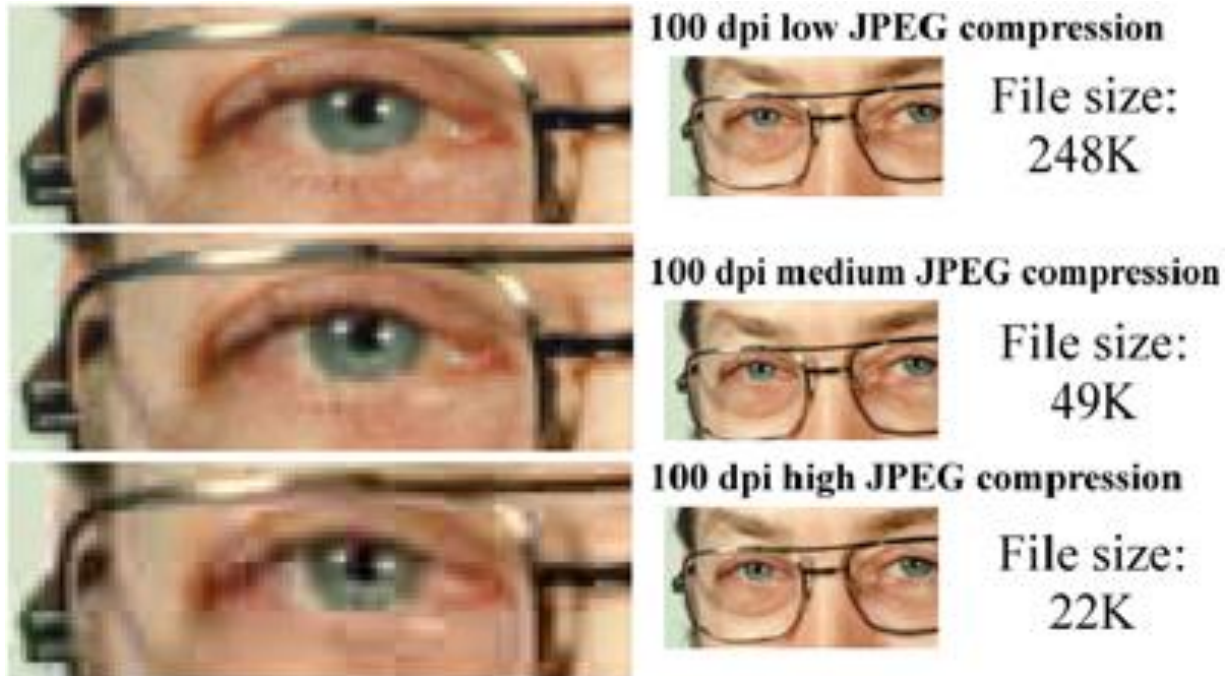
# A Complete System: JPEG

- A set of methods with a common *baseline* system
  - Discrete Cosine Transform
  - Quantisation
  - Variable length encoding
- A JPEG-compatible product must only include support for the baseline



# JPEG Compression

- Increasing the amount of quantisation reduces file size but introduces artefacts: blocks become visible



# NEXT WEEK

---

CNN

Conclusion and Exam Revision