

COMP2054-ADE

Algorithms Data Structures & Efficiency

ADE Lec01

Analysis of Algorithms


Lecturer: Andrew Parkes

Email: andrew.parkes@nottingham.ac.uk

<http://www.cs.nott.ac.uk/~pszajp/>

Running Time: “finite” but how big?


- Consider “batch algorithms”:
transform input data into output data – as opposed to “interactive”
- The running time of an algorithm typically grows with the input size.
- Even at given size the runtime is usually not fixed.
 - So have “best”, “average” and “worst” cases.
 - A typical example →
- We (usually) focus on the worst case running time at given size
 - Useful, and easier to analyse
 - Average case time is often difficult to determine.



Experimental Studies

General Pattern:

- Write a program implementing the algorithm
- Run the program with inputs of “varying size and composition”
- Use a system method to get an (in)accurate measure of the actual running time
- Plot the results
 - Example is shown →
- Interpret & analyse. E.g. is it
 - A “power law”, n^k , for some k
 - An “exponential”, b^n , for some b



Limitations of Experiments

- It is necessary to implement the algorithm, which may be difficult or time-consuming
- **Results may not be indicative of the running time on other inputs not included in the experiment.**
 - **Maybe we miss the “real worst case”**
- In order to compare two algorithms directly, the same hardware and software environments must be used

Limitations of Theory

- It is necessary to implement the theory, which may be difficult or time-consuming
- Results may not be indicative of the typical running time on inputs encountered in real world.

So can be useful to be able to use both experiment and theory.

Aside: Theory vs. Experiment

Standard science:

**“Never believe a theory until it has been
‘confirmed’ by an experiment”**

Partially joking:

**“Never believe an experiment until it has
been confirmed by a theory”**

- **Attributed to Sir Arthur Eddington**
<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3597502/>

Theoretical Analysis

- **AIM: Characterise running time as a function of the input size, n .**
- Uses a “high-level” description of the algorithm instead of an implementation
 - Takes into account all possible inputs
 - Allows us to evaluate the speed of an algorithm independently of the hardware/software/language environment

Pseudocode (recap)

- High-level description of an algorithm
- More structured than English prose
- Less detailed than a program
- Preferred notation for describing algorithms
- Hides program design issues

Example: find max element of an array

Algorithm *arrayMax*(*A*, *n*)

Input array *A* of *n* integers

Output maximum element of *A*

currentMax $\leftarrow A[0]$

for *i* $\leftarrow 1$ **to** *n* $- 1$ **do**

if *A*[*i*] > *currentMax* **then**

currentMax $\leftarrow A[i]$

return *currentMax*

Pseudocode Details (recap)

- Control flow
 - **if** ... **then** ... [**else** ...]
 - **while** ... **do** ...
 - **repeat** ... **until** ...
 - **for** ... **do** ...
 - Indentation replaces braces
- Method declaration

Algorithm *method* (*arg* [, *arg*...])

Input ...

Output ...
- Method call


var.method (*arg* [, *arg*...])
- Return value

return *expression*
- Expressions
 - ← Assignment
(like = in Java)
 - = Equality testing
(like == in Java)
 - n^2 Superscripts and other
mathematical
formatting allowed

Primitive Operations

- Basic computations performed by an algorithm
 - Identifiable in pseudocode
 - Largely independent from the programming language
 - Exact definition not important (we will see why later)
- Assumed to take a constant amount of time in the “RAM model” (next slide)
 - Tend to be close to “Assembly language”
 - No “hidden expenses”
- Examples:
 - Assigning a value to a variable
 - Indexing into an array
 - Comparing two numbers
 - Adding/subtracting/multiplying/dividing two numbers
 - Calling a method
 - Returning from a method

The Random Access Machine (RAM) Model



- A **CPU**
- A potentially-unbounded bank of **memory** cells, each of which can hold an arbitrary number or character
- Memory cells are numbered and accessing any cell in memory takes unit time (some fixed time).
 - (Note that RAM can stand for both “Random Access Machine” and “Random Access Memory,” Which is an unfortunate, but standard, over-loading of terminology.)

Limitations of RAM model

- “... can hold an arbitrary number ...” ?
 - Can we really expect to store
“938566359286151800351666177735777777177177374717717471
5717777761365661618161616” in one cell on a real computer?
- Here, we ignore such “bignum” issues. Instead:
- “all numbers are of equal size, as they all fit in a single register of the CPU”
- 64bits (signed int) allows up to 9,223,372,036,854,775,807
 - Exercise (offline): compare to: nanoseconds since big bang;
national debt.
- Note: on real machines (usually) computing
1 + 1 takes as long as 381513 + 243542
Hence, we typically ignore the sizes of numbers in the arithmetic operations.

Counting Primitive Operations

- By inspecting the pseudocode, we can determine the maximum number of primitive operations executed by an algorithm, as a function of the input size

Algorithm <i>arrayMax</i> (<i>A</i> , <i>n</i>)	# operations
<i>currentMax</i> $\leftarrow A[0]$?
for <i>i</i> $\leftarrow 1$ to <i>n</i> - 1 do	?
if <i>A</i> [<i>i</i>] > <i>currentMax</i> then	?
<i>currentMax</i> $\leftarrow A[i]$?
return <i>currentMax</i>	?
Total	<u>EXERCISE “Try it!”</u>

Counting Primitive Operations (partial) – “Offline -> Pause”

- **Worst case** number of primitive operations executed as a function of the input size, n

Algorithm <i>arrayMax</i> (A, n)	# operations
<i>currentMax</i> $\leftarrow A[0]$?
for $i \leftarrow 1$ to $n - 1$ do	?
if $A[i] > \textit{currentMax}$ then	?
<i>currentMax</i> $\leftarrow A[i]$	$2(n - 1)$
	?
return <i>currentMax</i>	?
Total	??

Counting Primitive Operations (partial)

- Worst case number of primitive operations executed as a function of the input size, n

Algorithm <i>arrayMax</i> (A, n)	# operations
<i>currentMax</i> $\leftarrow A[0]$?
for $i \leftarrow 1$ to $n - 1$ do	1 // for $i \leftarrow 1$
if $A[i] > \textit{currentMax}$ then	2 ($n - 1$)
<i>currentMax</i> $\leftarrow A[i]$??
 return <i>currentMax</i>	 ?
Total	??

Counting Primitive Operations (all)

- Worst case number of primitive operations executed as a function of the input size, n

Algorithm <i>arrayMax</i> (A, n)	# operations
<i>currentMax</i> $\leftarrow A[0]$	2
for $i \leftarrow 1$ to $n - 1$ do	1
if $A[i] > \textit{currentMax}$ then	$2(n - 1)$
<i>currentMax</i> $\leftarrow A[i]$	$2(n - 1)$ (worst case)
{ increment counter: $i++$ }	$2(n - 1)$ (“hidden”)
{ test counter: $i \leq (n-1)$ }	$2(n - 1)$ (“hidden”)
return <i>currentMax</i>	1
Total	$8n - 4$

Counting is “underspecified”

- Consider “ $c \leftarrow A[i]$ ” then ‘full’ process can be
 - get A = pointer to start of array A , and store into a register
 - get i , and store into a register
 - compute $A+i$ = pointer to location of $A[i]$, and store back into a register
 - get value of “ $*(A+i)$ ” (from RAM) and store value it into a register
 - copy the value into the location of c in the RAM
- might not want to count all this, e.g. just count
 - ‘plus’ of “ $A+i$ ”
 - the assignment

Counting is “underspecified”

- There can be multiple right answers – if you get ‘2’ and I count ‘4’ then it does not mean you are wrong!
- Note: If I think an answer is ‘4’ then ‘2’ is probably also acceptable – but “2 n” probably will not be.
- It is most important to be able to
 - know what is happening in the underlying process
 - be able to link to C and assembly level notions
 - be able to use this to give a reasonably consistent justification of your answers

Note: Correctness vs. Efficiency

- Primitive operation counting is relevant to Efficiency – but not (directly) for Correctness
- For correctness, do not about runtime of the alg.
 - Do care about the time to find a proof if doing an automated search for proofs.
 - The verification seems to be quick in lean: e.g.
`#eval 2035713999 + 350135299`
 - Very quick (<1sec) – not using succ internally (?)
- Note that we did not prove the algorithm correct
 - Would need to do arrays/lists in Lean first
 - Just for thought: And then what?
 - Can do efficiency of incorrect algorithms 😊

Estimating Running Time

- Algorithm *arrayMax* executes $8n - 4$ primitive operations in the worst case. Define:
 - a = Time taken by the fastest primitive operation
 - b = Time taken by the slowest primitive operation
- Let $T(n)$ be worst-case time of *arrayMax*. Then
$$a(8n - 4) \leq T(n) \leq b(8n - 4)$$
- Hence, $T(n)$ is bounded 'above and below' by two linear functions
- Usually said as "*arrayMax* runs in linear time"

Remarks

- Do not get too obsessed with the fine details of counting of primitive operations
- The details of the counting and timing would probably depend
 - the compiler, and require inspection of the assembly code
 - the CPU architecture, pipelining, cache misses, etc, etc

Growth Rate of Running Time

- Changing the hardware/ software environment
 - Affects $T(n)$ by a constant factor, but
 - Does not alter the growth rate of $T(n)$
- **The linear growth rate of the running time $T(n)$ is an intrinsic property of algorithm *arrayMax***

Exercise: (exam-style question)

Given the following code fragment:

```
m ← 0  
while (n ≥ 2)  
    n ← n/2  
    m++  
return m
```

Give an analysis of its runtime

Exercise: what is $T(n)$ of alg-lec1?

Algorithm: *alg-lec1*

Input: positive integer n , which is a power of 2.

I.e. there exists k such that $2^k = n$

$m \leftarrow 0$

***while* ($n \geq 2$)**

$n \leftarrow n/2$

$m++$

***return* m**

Exercise: what is $T(n)$ of alg-lec1? (cont)

Algorithm: *alg-lec1*

Input: positive integer n , which is a power of 2

Output: integer m such that $2^m = n$

$m \leftarrow 0$	1
<i>while</i> ($n \geq 2$)	?
$n \leftarrow n/2$?
$m++$?
<i>return</i> m	1

Exercise

Algorithm: *alg-lec1*

Input: positive integer n , which is a power of 2

Output: integer m such that $2^m = n$

$m \leftarrow 0$

1

***while* ($n \geq 2$)**

? per pass

$n \leftarrow n/2$

? per pass

$m++$

? per pass

***return* m**

1

(Pause and try/think)

Internal Steps:

$$n \leftarrow n/2 \quad ?$$

1. read ***n*** from memory (RAM) and store in a register r1 (very fast piece of memory on the CPU)
2. read **2** from memory and store in a register r2
3. send registers r1 r2 through arithmetic division and store result in a register r3
4. write r3 back to ***n***

CPU steps needed is 4, does not depend on ***n***

Internal Steps: different compiler

$$n \leftarrow n/2$$

?

1. read ***n*** from memory (RAM) and store in a register r1 (very fast piece of memory on the CPU)
2. send registers r1 through a right shift of the bits and store result in a register r3
e.g. compute $13/2=6$ by $1101 \rightarrow 110$
3. write r3 back to ***n***

CPU steps needed is 3, different but still does not depend on ***n***

Exercise (cont)

Algorithm: *alg-lec1*

Input: positive integer n , which is a power of 2

Output: integer m such that $2^m = n$

$m \leftarrow 0$

1

$\text{while } (n \geq 2)$

3 per pass

$n \leftarrow n/2$

3 per pass

$m++$

3 per pass

$\text{return } m$

1

Thought Exercise (offline)

- Based on your knowledge of assembly, and machine architectures try to estimate the number of CPU cycles that might actually be used.
- “Divide” or “shift”? Which is faster?
- **Point: try to eventually build a mental model that is an “internal interpreter” so as to know how a program will run**
- **Such “internal interpreters” are vital for understanding programming (IMHO)**

How many passes through the loop of alg-lec1?

Hint: If ever stuck:

- Try simple concrete examples
- Start from “ridiculously simple” and work up to harder examples

(Real mathematicians often work this way; but then hide it when writing up 😞)

- Do a “trace of the program” by hand:

How many passes through loop?

- Focus on the relevant portions:

while ($n \geq 2$)

$n \leftarrow n/2$

- Simplest example?
- Exercise: What is smallest positive integer that is a power of two?
- Answer: 1 as $2^0 = 1$
- (If confused, or if you answered "2", then consider revising your maths about exponents and logarithms)

How many passes through loop? (cont)

while ($n \geq 2$) { $n \leftarrow n/2$; }

- Case: $n = 1 = 2^0$ passes = 0
- Case: $n = 2 = 2^1$
 - $n=2$, then $n=1$; passes=1
- Case: $n = 4 = 2^2$
 - $n=4$, then $n=2$, then $n=1$; passes = 2
- Case: $n = 8 = 2^3$
 - $n=8, 4, 2$, then $n=1$; passes = 3

How many passes through loop? (cont)

- Case: $n = 2^m$
 - $n=2^m, 2^{m-1}, 2^{m-2}, \dots, 2$
 - m passes through loop
- but note, n is the input not m , so want to write answer in terms of n . Use
 - $m = \log_2(n)$
- Result: passes through loop = $\log_2(n)$

Exercise (cont)

Algorithm: *alg-lec1*

Input: positive integer n , which is a power of 2

Output: integer m such that $2^m = n$

$m \leftarrow 0$

1

$\text{while } (n \geq 2)$

3 $(\log_2(n)+1)$

$n \leftarrow n/2$

3 $\log_2(n)$

$m++$

3 $\log_2(n)$

$\text{return } m$

1

all together: $9 \log_2(n) + 5$

(the "+1" on line 2, is because the test is done even if it fails)

****Remarks****

- Each pass through the loop the size of n is halved
 - the “ $\log_2(n)$ ” is typical of such “halving on each iteration”
- **This concept also appears in sorting and searching; hence you MUST make sure you fully understand this example**
 - **The ADE half of the module will probably be incomprehensible otherwise**

Summary

Goal:

- Build foundations for time-analysis of programs

Skills needed:

- Count primitive operations
- Counting of operations with
 - Loops
 - (Recursion)

Removing details

- According to precisely how we count steps we might get many different answers, e.g. something like
 - $5 \log_2(n) + 2$
 - $9 \log_2(n) + 5$, etc
- Also this counts “steps”
 - the translation to runtime depends on the compiler, hardware, etc
- **Need a way to suppress such details**

Next Lecture

“Suppressing the details”

A motivation and introduction to big-Oh

Oh... ohh!!

Exercise: (Advanced, Offline/Take-home)

Given the following code fragment:

```
m ← 0  
while (n ≥ 2)  
    n ← sqrt(n)  
    m++  
return m
```

Give an analysis of its runtime.

For the counting, assume that the square root "sqrt" is a primitive operation.