

COMP2054-ADE

Map ADT and hashtables

The Map ADT over $\langle K, V \rangle$

Map ADT methods:

- **V** **get**(K k): if the map M has an entry with key k, return its associated value; else, return null
- **V** **put**(K k, V v): insert entry (k, v) into the map M; if key k is not already in M, then return null; else, return old value associated with k
- **V** **remove**(K k): if the map M has an entry with key k, remove it from M and return its associated value; else, return null
- **int** **size**(), **boolean** **isEmpty**()
- **{K}** **keys**(): return an iterable collection of the keys in M
- **{V}** **values**(): return an iterable collection of the values in M
- **{ $\langle K, V \rangle$ }** **entries**(): return an iterable collection of the entries in M

Hash Tables

- **Hash tables** are a concrete data structure which is suitable for implementing maps.
- **Basic idea: convert each key into an index into a (big) array.**
- Look-up of keys and insertion and deletion in a hash table usually runs in $O(1)$ time.
 - Not guaranteed, and design of the table needs to be done carefully if want the access to be “reliably $O(1)$ ”


Hash Functions and Hash Tables

- A **hash function** h maps keys of a given type to integers in a fixed interval $[0, N - 1]$
 - Example:
$$h(k) = k \bmod N$$

is a hash function for integer keys
- The integer $h(k)$ is called the **hash value** of key k
- A **hash table** for a given key type consists of
 - Hash function h
 - Array (called table) of size N
- When implementing a map with a hash table, the goal is to store item (k, v) at index $i = h(k)$


Example

- We design a hash table for a map storing entries as (SID, Name), where SID (student identity number) is an eight-digit positive integer
- Our hash table uses an array of size $N = 10,000$ and the hash function
$$h(x) = \text{last four digits of } x$$
$$= "x \bmod 10000"$$
(details depends if SID is stored as an int or a string)



Collision Handling

- Collisions occur when different elements are mapped to the same cell
- A lot of the theory and practice of hashing consists of devising better ways to avoid or handle collisions



Hash Functions

- A hash function is usually specified as the composition of two functions:

Hash code:

$h_1: \text{keys} \rightarrow \text{integers}$

Compression function:

$h_2: \text{integers} \rightarrow [0, N - 1]$

- The hash code is applied first, and the compression function is applied next on the result, i.e.,
$$h(x) = h_2(h_1(x))$$
- The goal of the hash function is to “disperse” the keys in an “apparently random” way

Hash functions: “dispersal”?

- Re: The goal of the hash function is to “disperse” the keys in an “apparently random” way
- Questions:
 - Why disperse?
 - Why random?

Hash functions: “dispersal”?

- Why disperse?
 - to reduce numbers of collisions
- Why random?
 - random means ‘no pattern’
 - if there is an obvious pattern then the incoming data might have a matching pattern that leads to many collisions
 - “sometimes ‘no pattern’ is the only safe pattern” (e.g. rock-paper-scissors game)

Hash Codes [Not assessed]

- **Memory address:**
 - We reinterpret the memory address of the key object as an integer (default hash code of all Java objects)
 - Good in general, except for numeric and string keys
- **Integer cast:**
 - We reinterpret the bits of the key as an integer
 - Suitable for keys of length less than or equal to the number of bits of the integer type (e.g., byte, short, int and float in Java)
- **Component sum:**
 - We partition the bits of the key into components of fixed length (e.g., 16 or 32 bits) and we sum the components (ignoring overflows)
 - Suitable for numeric keys of fixed length greater than or equal to the number of bits of the integer type (e.g., long and double in Java)

Hash Codes (cont.) [Not assessed]

- **Polynomial accumulation:**

- We partition the bits of the key into a sequence of components of fixed length (e.g., 8, 16 or 32 bits)

$$a_0 a_1 \dots a_{n-1}$$

- We evaluate the polynomial

$$p(z) = a_0 + a_1 z + a_2 z^2 + \dots \\ \dots + a_{n-1} z^{n-1}$$

at a fixed value z , ignoring overflows

- Especially suitable for strings (e.g., the choice $z = 33$ gives at most 6 collisions on a set of 50,000 English words)

- Or, representing a string as a number on base z

- Compare with base 10:

$$365 = 3 \cdot 10^2 + 6 \cdot 10^1 + 5 \cdot 10^0$$

- Base 27 (26 characters + blank):

$$\text{cab} = 3 \cdot 27^2 + 1 \cdot 27^1 + 2 \cdot 27^0$$

where

$$a = 1, b = 2, c = 3 \text{ and } z = 26$$

Compression Functions

- Division:

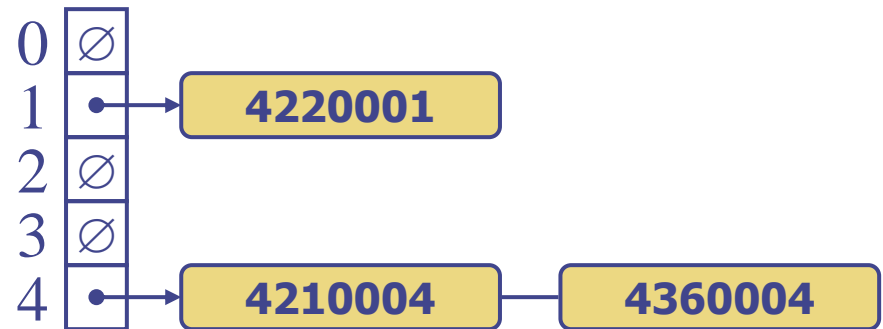
- $h_2(y) = y \bmod N$
- The size N of the hash table is usually chosen to be a prime (hash codes will tend to spread better)

- Multiply, Add and Divide (MAD):

- $h_2(y) = (ay + b) \bmod N$
- a and b are nonnegative integers
- such that
 - $a \bmod N \neq 0$
 - Otherwise, every integer would map to the same value b

Collision Handling

- Collisions occur when different elements are mapped to the same cell
- **Separate Chaining:** let each cell in the table point to (e.g.) a linked list of entries that map there
 - Note: In practice, should use a more efficient Map; e.g. a Binary Search Tree (BST), (see later lectures)



Map Methods with Separate Chaining used for Collisions

- Delegate operations to a list-based map at each cell:

Algorithm $\text{get}(k)$:

Output: The value associated with the key k in the map, or **null** if there is no entry with key equal to k in the map

return $A[h(k)].\text{get}(k)$

// Simply delegates the “get” to the list-based map at $A[h(k)]$

Map Methods with Separate Chaining used for Collisions

Algorithm $\text{put}(k, v)$:

Output: If there is an existing entry in our map with key equal to k , then we return its value (replacing it with v); otherwise, we return **null**

$t \leftarrow A[h(k)].\text{put}(k, v)$

// Simply delegates the put to the list-based map at $A[h(k)]$

if $t = \text{null}$ **then**

$\{k \text{ is a new key}\}$

$n \leftarrow n + 1$

return t

Map Methods with Separate Chaining used for Collisions

Algorithm `remove(k)`:

Output: The (removed) value associated with key k in the map, or **null** if there is no entry with key equal to k in the map

$t \leftarrow A[h(k)].remove(k)$

// Simply delegates the remove to the list-based map at $A[h(k)]$
if $t \neq \text{null}$ **then** { k was found}

$n \leftarrow n - 1$

return t

*Access is still $O(n)$, but usually the relevant n is much smaller, because it usually builds many small lists, e.g. length n/N on average.
So this method might be okay, sometimes.*

Separate Chaining

- Separate chaining is simple and fast, but requires additional memory outside the table.
- When memory is critical then we try harder to remain within the existing memory:

Open addressing

- **Open addressing**: the colliding item is placed in a different cell of the table
- **Linear probing** handles collisions by placing the colliding item in the next (circularly) available table cell (variant: $\text{cell} + c$ where c is a constant)
 - “Circular array” – once get to the right-hand end, then just start again at the beginning of the array
- Each table cell inspected is referred to as a “probe”
- Disadvantage: Colliding items lump together, causing future collisions to cause a longer sequence of probes

Open addressing : Example

- Example: $h(x) = x \bmod 13$, $c=1$
 - Insert keys 18, 41, 22, 44, 59, 32, 5 in this order
 - 18,41,22 have no collisions, giving

		41			18				22			
--	--	----	--	--	----	--	--	--	----	--	--	--

- But $44 \bmod 13 = 5$ collides with $18 \bmod 13 = 5$, so use index $5+1$

		41			18	44			22			
--	--	----	--	--	----	----	--	--	----	--	--	--

- Now $59 \bmod 13 = 7$ has no collisions

		41			18	44	59		22			
--	--	----	--	--	----	----	----	--	----	--	--	--


- Now $32 \bmod 13 = 6$ collides with the 44, and so walk past the 44 and 59 to find an empty space

		41			18	44	59	32	22			
--	--	----	--	--	----	----	----	----	----	--	--	--

- Now $5 \bmod 13 = 5$ has a long walk to find an empty space:

		41			18	44	59	32	22	5		
--	--	----	--	--	----	----	----	----	----	---	--	--

Search with Linear Probing



- Consider a hash table A that uses linear probing
- **get(k)**
 - We start at cell $h(k)$
 - We probe consecutive locations until one of the following occurs
 - An item with key k is found, or
 - An empty cell is found, or
 - N cells have been unsuccessfully probed

```
Algorithm get(k)  
   $i \leftarrow h(k)$   
   $p \leftarrow 0$   
  repeat  
     $c \leftarrow A[i]$   
    if  $c = \emptyset$   
      return null  
    else if  $c.key() = k$   
      return  $c.element()$   
    else  
       $i \leftarrow (i + 1) \bmod N$   
       $p \leftarrow p + 1$   
  until  $p = N$   
  return null
```

Open addressing : Example

- Example: $h(x) = x \bmod 13$, $c=1$
 - Insert keys 18, 41, 22, 44, 59, 32, 5 in this order
 - Gave

		41			18	44	59	32	22	5		
--	--	----	--	--	----	----	----	----	----	---	--	--

- Suppose we now remove(22).
It seems we should just simply obtain:

		41			18	44	59	32		5		
--	--	----	--	--	----	----	----	----	--	---	--	--

- Do you see a problem?
- Suppose we try to do get(5)
- get(5) would fail to find the 5
- as the scan stops at the empty cell where 22 used to be

Exercise

- How do you safely remove an element x ?
- One answer:
 - (not the best answer, but simplest):
 - Find x using ``get'` and set the entry back to blank, i.e. null or empty (which sometimes write as ``#'`)
 - Fix the sequence on its right-hand-side
 - WHY!?: If any entry on the right used linear probing then it might no longer be discoverable by ``get'` because it will stop at the blank!!!
 - Fix: Move such entries, e.g. by removing them and then re-inserting them all.
 - **EXERCISE (offline)** Figure out the details and write pseudo-code for this and do examples. (Ask in labs/tutorials if needed!)

Exercise

How do you safely remove an element x ?

Another solution:

- “**Lazy deletion**”: don’t mark the entry as a blank, but as a ‘deleted’ and fix the entries later. E.g. see

http://opendatastructures.org/versions/edition-0.1e/ods-java/5_2_LinearHashTable_Linear.html

- E.g. in the find, skip over a ‘deleted’ entry rather than stopping

(See Tutorials)

Double Hashing

- Double hashing uses a secondary hash function $d(k)$ and handles collisions by placing an item in the first available cell of the series
$$(h(k) + j d(k)) \bmod N$$
for $j = 0, 1, \dots, N - 1$
- The secondary hash function $d(k)$ cannot have zero values
- Linear probing is just $d(k)=1$
- The table size N must be a prime to allow probing of all the cells
- Common choice for the secondary hash function:
$$d(k) = q - (k \bmod q)$$
where
 - $q < N$
 - q is a prime
- The possible values for $d(k)$ are
$$1, 2, \dots, q$$


Remarks

- “The table size N must be a prime to allow probing of all the cells”
- E.g. consider $d(k) = 4$
 - With $N=12$ then the only positions scanned are 4,8,0,4, ...
 - So we miss many cells that might have space
 - With $N=11$ then the positions are 4,8,1,5,9,2, ...
- With a prime N , then eventually all table positions will be probed

Example of Double Hashing

- Consider a hash table storing integer keys that handles collision with double hashing
 - $N = 13$
 - $h(k) = k \bmod 13$
 - $d(k) = 7 - (k \bmod 7)$
- Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order

k	$h(k)$	$d(k)$	Probes	
18	5	3	5	
41	2	1	2	
22	9	6	9	
44	5	5	5	10
59	7	4	7	
32	6	3	6	
31	5	4	5	9 0
73	8	4	8	



Example of Double Hashing

- Consider a hash table storing integer keys that handles collision with double hashing

- $N = 13$
- $h(k) = k \bmod 13$
- $d(k) = 7 - (k \bmod 7)$

- Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order

- Exercise: do the details yourself, and see tutorials

k	$h(k)$	$d(k)$	Probes	
18	5	3	5	
41	2	1	2	
22	9	6	9	
44	5	5	5	10
59	7	4	7	
32	6	3	6	
31	5	4	5	9 0
73	8	4	8	

0	1	2	3	4	5	6	7	8	9	10	11	12



31		41			18	32	59	73	22	44		
0	1	2	3	4	5	6	7	8	9	10	11	12

Performance of Hashing

- In the worst case, searches, insertions and removals on a hash table take $O(n)$ time
- The worst case occurs when all the keys inserted into the map collide
- The load factor $\alpha = n/N$ affects the performance of a hash table
- In Java, maximal load factor is 0.75 (75%) – after that, rehashed
 - as for Vector, it may be good to “roughly double” the table size each rehash
 - pick a new (prime) close to twice the current size
- The expected running time of all the map ADT operations in a hash table is $O(1)$
- In practice, hashing is very fast provided the load factor is not close to 100%

Re-Hashing

When the table gets too full then “re-hash”:
Create a new larger table and new hash function.

- Need to (eventually) transfer all the entries from the old table to the new one
- If do so immediately, then
 - one can amortise the cost over many entries (as for Vector) and so get an average cost of $O(1)$ again
 - but the worst case might be $O(n)$ when the table is rehashed, and this might be bad for a real time system
 - Option:
do not transfer all entries “in one go” but do “a few at a time”
 - Keep both tables until the transfer is complete; but only do insertions into the new table.

Exercise (offline): consider this in more detail, and read the relevant part of the text book (Section 9.2.7 Load factors and Reshasing) or the wiki page http://en.wikipedia.org/wiki/Hash_table#Dynamic_resizing

Applications of Hashing

- Direct applications of hash tables:
 - small databases
 - compilers
 - browser caches – the weird and wonderful filenames in the browser cache folder are hashcodes of something?
- Hash tables as an auxiliary data structure in a program:
 - Look-up table: if you want to check whether some object has been seen before, for example in a graph or list traversal, keep a hashtable of (object, "seen before") pairs, where the key is the reference to the object, and the value is some arbitrary marker.

Comparison of HashMap and PQ

- HashMap does not use the ordering of keys
 - E.g. does not implement `min()`
 - In a hash table it would need a scan of all the keys in the table, so $O(n)$ (or worse)
- PQ does not allow direct access to a key
 - E.g. there is no easy way to do `get(k)`
 - In a (standard) heap we would have to walk through all the entries

Minimum Expectations

- Map ADT, and its usage
- basic concepts of hash tables
- hash codes and compression functions
- Options to handle collisions
 - Separate chaining, linear probing, double hashing
 - How to insert, find/get, remove for all these systems