

COMP2054 ADE

Binary Search Trees

Note to help with understanding: “Global vs. Local Perspectives”

(intended to help with coding/understanding algorithms)

- Global view of tree (or any other data structure)
 - can look at entire tree in one go
 - human seeing a picture of a (small) tree
- Local view of tree
 - what your code sees: “code perspective”
 - code generally only sees a local portion of a tree and must work with that
- Vital coding skill: develop ability to see the local view, “code perspective”, and not only the global
 - at each code line need to know which data is immediately accessible
 - Avoid trying to code in a way that cannot be done locally
 - Note the local view still needs to work towards a global goal

Note intended to help with your coding: “Global vs. Local Perspectives” Analogies

- Your Data Structure:
 - A real tree
- Your code:
 - “A short-sighted ant crawling over the tree”

Note intended to help with your coding: “Global vs. Local Perspectives” Analogies

- Your Data Structure:
 - A big underground cave system
- Your code:
 - You!
 - ... with a weak flashlight
 - ... and no overall map
 - ... but a goal of which cavern to reach

Motivations

- Suppose you have an array A (of ints) and want to search for a particular int k
- If the array is unsorted then no choice but to scan all elements, hence $O(n)$
- If it is sorted then we can do a lot better. How?

Motivations

- Binary Search for an element within a sorted array:

Algorithm *BinarySearch*(int[] A, int *k*, int *L*, int *R*)

// searches for k in the range A[L] ... A[R] of the sorted array A

// (note the standard technique of working with a sub-array)

if $R < L$

 return false

$m = \text{floor}((L + R) / 2)$ // an approximate midpoint

if $k == A[m]$

 return true

else if $k < A[m]$

 return *BinarySearch*(A, k, L, m-1)

else // $k > A[m]$

 return *BinarySearch*(A, k, m+1, R)

Motivations

- Binary Search for an element within a sorted array is fast
 - The array to be searched is halved at each iteration
 - Hence $O(\log n)$ (Next slide proves this).
- It only works because of the step of “knowing whether to go left or right”
- But arrays suffer from being slow, $O(n)$, to insert new elements
 - because need to shift $O(n)$ elements to make room for them
- **Search trees attempt to fix the inefficiency of insertion whilst keeping good $O(\log n)$ properties of binary search**
 - **So, we need a tree where we “know which way to go”:**

Binary search complexity

- Using recurrence relations
 - $T(n) = T(n/2) + 1$
 - “ $n/2$ ” because only have to search $1/2$ the array after recursion
 - “1” as the cost of doing the comparison and deciding whether to go left or right
 - Can do exactly (exercise!)
- Use Master Theorem (revise lectures if needed):
 - $a=1, b=2, \log_b(a) = \log_2(1) = 0$
 - So $c = \log_b(a) = 0$
 - Have $f(n) = 1$ which is $\Theta(n^0)$
 - So is Case 2: and then is n^0 but “get an extra log”
 - Hence, $\Theta(\log(n))$

**** Binary Search Trees ****

- A binary search tree is a binary tree storing key-value entries at its internal nodes and satisfying the following “search tree” property:


- Let u , v , and w be any three nodes such that u is in the **left subtree** of v and w is in the **right subtree** of v .

Then we must have

$$\text{key}(u) \leq \text{key}(v) \leq \text{key}(w)$$


or, as we will assume there are no duplicate keys:

$$\text{key}(u) < \text{key}(v) < \text{key}(w)$$




Binary Search Trees

- “Null” children
 - Sometimes, just for clarity, we explicitly show the “null pointers” that are for the case of a non-existent child




Not a Binary Search Tree

- VERY IMPORTANT:
- The following is **not** a Binary Search Tree
 - Even though the immediate children of every node look correct, we need that **every** node in the left subtree of 6 is less than 6, and this is failed by the node '7'



Search

- To search for key k , trace a downward path starting at the root
- The next node visited depends on the outcome of the comparison of k with the key of the current node
- If we reach a leaf, the key is not found and we return null



Search

Algorithm *Node TreeSearch*(Key k , Node n)

if $n.isExternal()$ // or, “if $n == null$ ”

return *null*

if $k < n.key()$

return *TreeSearch*(k , $n.left()$)


else if $k = n.key()$

return n

else // $k > n.key()$


return *TreeSearch*(k , $n.right()$)

- Example: find(4):
 - Call *TreeSearch*(4, root)
 - Note: standard trick of using an auxiliary function that keeps more of the “state”
- Compare this with the binary search of an array




Search from “code perspective”

- Example: `find(4)`:
- $4 < 6$ so need to go left




Search from “code perspective”

- Example: `find(4)`:
- $4 > 2$ so need to go right



Search from “code perspective”

- Example: `find(4)`:
- Found 4, but note that we only saw/visited a small part of the entire tree




- It only follows a path from the root to some node, and so is clearly $O(h)$ where h is the height of the tree.

Recursive vs. Iterative

- Generally: Recursive programs are easier to implement, but less efficient
 - because of the overhead of a function call
- For best efficiency, will need to convert to an iterative program
 - “while (test) { ... }”
- Exercises (offline): convert to iterative
 - binary search of an array
 - search in a binary search tree

Fundamental Property of Search Tree

- Exercise: what does an inorder traversal of the following search tree produce?




Fundamental Property of Search Tree

Directly from the definition:

- An inorder traversal of a (binary) search trees visits the keys in increasing order

Fundamental Property of Search Tree

- Note that to access the minimum key, we just need to 'always go left'





Insertion


- Exercise: how to do insertion, `insert(k, v)` ?
- Have to insert `k` where a `get(k)` would find it!
- So natural that `insert(k,v)` starts with `get(k)`
- We search for key `k` (using `TreeSearch`)
- If `k` is already in the tree then just replace the value
- Otherwise, `k` is not already in the tree, and let `w` be the leaf reached by the search
 - We “insert `k` at node `w` and expand `w` into an internal node”
- Again, only follows a path from the root and so is $O(h)$

Insertion

- Example: insert 5
- We search for key $k=5$ (using TreeSearch)
- Let w be “missing child position” reached by the search
- We insert a real node, with key k , at that location w
- Exercise: check that this is correct! Does it preserve the search tree property?





Deletion



- How can we perform operation **remove(k)** ?
- As usual we start by trying to find(k)
- Four cases: (think of the externals as null – not as real children)
 1. k is not present, e.g. remove(7)
nothing to do
otherwise k is stored in some node n
 2. n has no children, e.g. remove(5),
(straightforward, left as an exercise)
 3. n has one child, e.g. remove(4)
 4. n has two children, e.g. remove(2)


Deletion – with one child

- Example: remove 4
- To perform operation `remove(4)`, we search for key 4.
- Let n be the node storing 4.
- Node n has a null left child, and a real child 5
- We remove n from the tree and connect 5 back to the parent of n
- **Note that this still works if n is a right child and has a left child. (Exercise).**



Deletion – with two children


- Example: remove 3
- **What can we do and still keep the property of being a search tree?**



- “Search tree” means an “inorder traversal” visits the keys “in order”
- As a sorted list, we would have [1,2,3,5,6,8,9]
- If we want to remove ‘3’ then a way to do this with minimal change to the rest of the list is to copy a key k' that is adjacent to ‘3’ on top of ‘3’ and then delete that key k'
 - Options in this case are k' being 2 or 5. We will focus on the nextKey, that is, ‘5’



Deletion – with two children

- Example: remove 3
- **What can we do and still keep the property of being a search tree?**



- “Search tree” means an “inorder traversal” visits the keys “in order”
- If these were a sorted list then “adjacent in the sequence” nodes 2 and 5 could be involved in the removal of node 3
- Can the next key after 3 ever have any left children?
- Can the next key after 3 ever have any children?

Deletion – with two children




- Can the next (using inorder) key after 3 ever have any children?
- YES. The right hand tree is also a possible search tree with same keys
- To access the nextKey, we need to go right, and then recursively-go-left
- IMPORTANT: that the next key from 3 cannot have a left child:
Proof by contradiction. If it did then the left child would not be “nextKey”

Deletion – with two children

Example: remove 3

The key node n has two internal children

1. we find the internal node w that follows n in an inorder traversal
2. we copy $key(w)$ into node n
3. we remove node w by means of same procedure as before for “one child”



Deletion – with two children

General procedure:

- Consider the case where the key k to be removed is stored at a node n whose children are both internal
 - we find the internal node w that follows n in an inorder traversal
 - we copy $key(w)$ into node n
 - we remove node w and its left child z (which must be a leaf) by using the same procedure as discussed earlier for “remove node with one child”
- Exercises (offline):
 - convince yourself this works properly even when w has a right child
 - Check that it is still $O(h)$


Exercises (offline)

In the “2 children deletion”, we had
“we find the internal node that follows
 n in an inorder traversal”

1. Could it also have been “precedes” instead of “follows”?
2. Might making a good choice of “follows” of “precedes” have some advantages?



Balanced Trees

- Binary search trees: if all levels filled, then search, insertion and deletion are $O(\log N)$.
 - As they are all $O(\text{height})$
- However, performance may deteriorate to linear if nodes are inserted in order:




Performance

- Consider a binary search tree of height h with n items
 - the space used is $O(n)$
 - methods **find**, **insert** and **remove** take $O(h)$ time
 - EXERCISE: carefully study the methods and check that this is true
- The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case




Issues in “Self-Balancing”

- Suppose you have a very imbalanced search tree – is there always a corresponding balanced search tree?
- Exercise (online): find a balanced tree for



Example



- For any number of nodes, there is (obviously?) a balanced tree
- Could balance by a “total rebuild”
 - just placing the keys into a new tree in the correct order

Solution: Self-Balancing

Goal of “Self-Balancing”:

- Constantly re-structure the trees:
- Keep the trees height balanced so that the height is logarithmic in the size
- Performance always logarithmic.

Issues in Self-Balancing

- Suppose you have a very imbalanced search tree, there are always corresponding balanced search trees
- Could make trees balanced using a “total rebuild”
 - But would require $O(n)$, and so very inefficient compared to the desired $O(\log n)$
- Re-balancing needs to be $O(\log n)$ or $O(\text{height})$
- Suggests re-balancing needs to just look at the path to some recently changed node, not the entire tree
- A priori, it is not at all obvious that this is possible!
- In a later lecture we show how to use “rotations” to adjust the tree whilst preserving the ordering

Minimal Expectations

- The definitions of Map ADT, BST, etc
- How to use binary search trees
- The complexity of operations as function of nodes n and height h