



University of  
**Nottingham**  
UK | CHINA | MALAYSIA

# COMP3055


# Machine Learning

## Topic 5 – Machine Learning Theory and Practice

Ying Weng  
2024 Autumn

# Motivating Example

## Polynomial Curve Fitting



Plot of a training data set of  $N = 10$  points, each comprising an observation of the input variable  $x$  along with the corresponding target variable  $t$ . The green curve shows the function  $\sin(2\pi x)$  used to generate the data. Our goal is to predict the value of  $t$  for some new value of  $x$ , **without** knowledge of the green curve.

# Classification VS Regression

## **Classification**

predict a label (discrete value)

## **Regression**

predict a response (continuous value)

# Regression Problem

- $X \equiv (x_1, \dots, x_N)^T$ ,  $T \equiv (t_1, \dots, t_N)^T$
- **Training set** is generated with  $t_n = \sin(2\pi x_n) + b_n$ ,  $n = 1, 2, \dots, N$ , where  $b_n$  is **random noise** having a Gaussian Distribution.
- **Goal**: predict the target value of  $t$  for some new input value  $x \rightarrow$  implicitly trying to discover the underlying function  $\sin(2\pi x_n)$ .
- Generalize from a finite data set ( $N=10$ )
- **Uncertainty**: the observed data are corrupted with noise.

# Polynomial Function

- We fit the training data using a **polynomial function** (linear models) of the form:

$$y(x, \mathbf{w}) = w_0 + w_1x^1 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

- Where  $M$  is the **order of the polynomial** (degree of freedom), and the **polynomial coefficients**  $w_0, w_1, \dots, w_M$  are collectively denoted by the vector  $\mathbf{w}$

# Solve a Polynomial Function


**Target:** find the values of *polynomial coefficients*

- Step 1: Fit the polynomial to the training data
- Step 2: minimize the *error function*

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

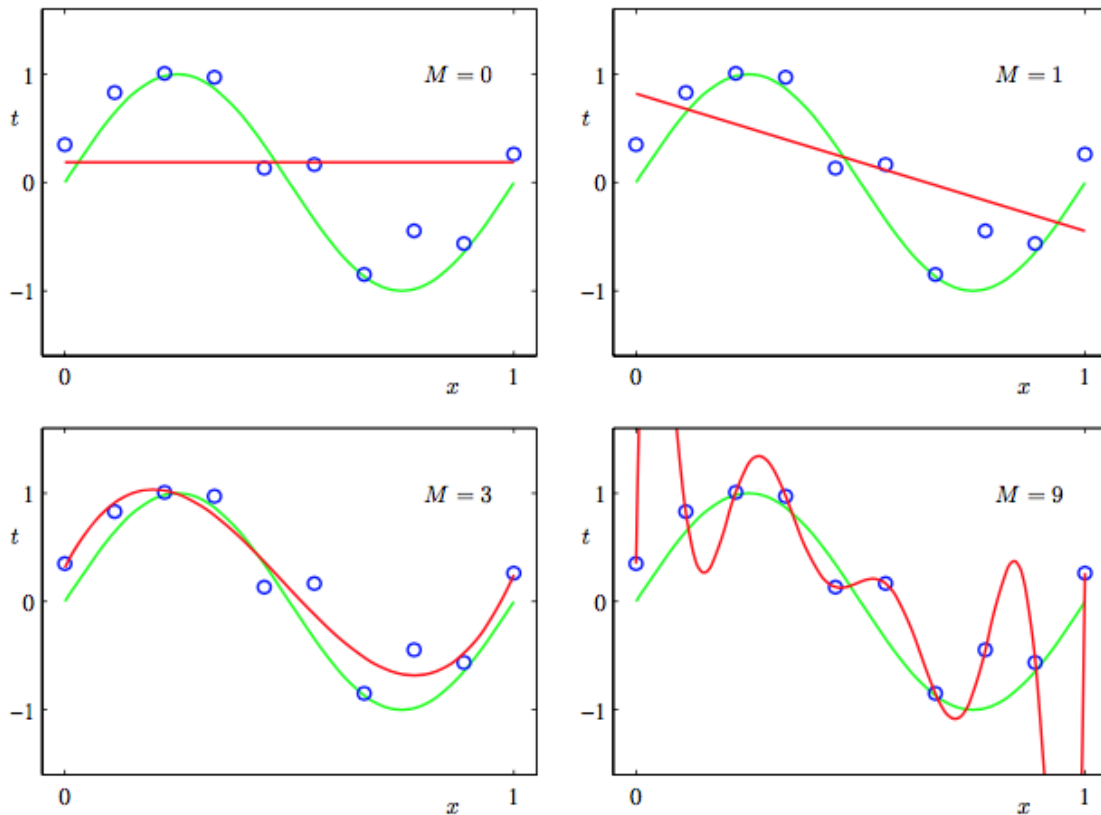
Where  $E(\mathbf{w})$  measures the misfit between the function  $y(x, \mathbf{w})$  and the training set data points

$\frac{1}{2}$  is included for later convenience.  $E(\mathbf{w})=0$  if and only if the function  $y(x, \mathbf{w})$  were to pass exactly through each training data point. The error function is a quadratic function of  $\mathbf{w}$ , there is only one solution.



# Model Selection

Choosing **order**  $M$  of the polynomial



# Model Selection

- **Goal:** achieve good *generalization* by making accurate predictions for new data.




- We use Root-mean-square (RMS) error on data

$$E_{RMS} = \sqrt{2E(w^*)/N}$$

- Where  $N$  allows us to compare different sizes of data set, and  $w^*$  is the *solution* of minimizing  $E(w)$  (*hypothesis*).
- It measures how well the model  $w^*$  doing in predicting the values of  $t$  for new data observations of  $x$ .



# Overfitting





	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

# Overfitting

Overfitting can occur when:

- Learning is performed for **too long** (e.g. in Neural Networks).
- The examples in the training set are **not representative** of all possible situations (is usually the case!).
- Model parameters are adjusted to **uninformative features** in the training set that have no causal relation to the true underlying target function!

# Overfitting



Increasing the size of the data set reduces the overfitting problem.

# Regularization


- *Regularization* can control the overfitting phenomenon, by adding *penalty term* to the error function to discourage the coefficients from reaching large values.

$$E(w) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, w) - t_n\}^2 + \frac{\lambda}{2} \|w\|^2$$

- Where  $\|w\|^2 \equiv w^T w = w_0^2 + w_1^2 + \dots + w_M^2$ , and the coefficient  $\lambda$  governs the relative importance of the regularization term compared with the sum-of-squares error term.

# Regularization


	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01



- Table of the coefficients  $w^*$  for  $M = 9$  polynomials with various values for the regularization parameter  $\lambda$ .
- $\lambda$  controls the *effective complexity* of the model and hence determines the *degree of overfitting*

# Cross Validation

- **Idea #1:** Choose hyperparameters that work best on the data



- **Idea #2:** Split data into **train** and **test**, choose hyperparameters that work best on test data



- **Idea #3:** Split data into **train**, **val**, and **test**; choose hyperparameters on **val** and evaluate on **test**



# Cross Validation

- **Idea #4: Cross-Validation:** Split data into **folds**, try each fold as validation and **average the results**

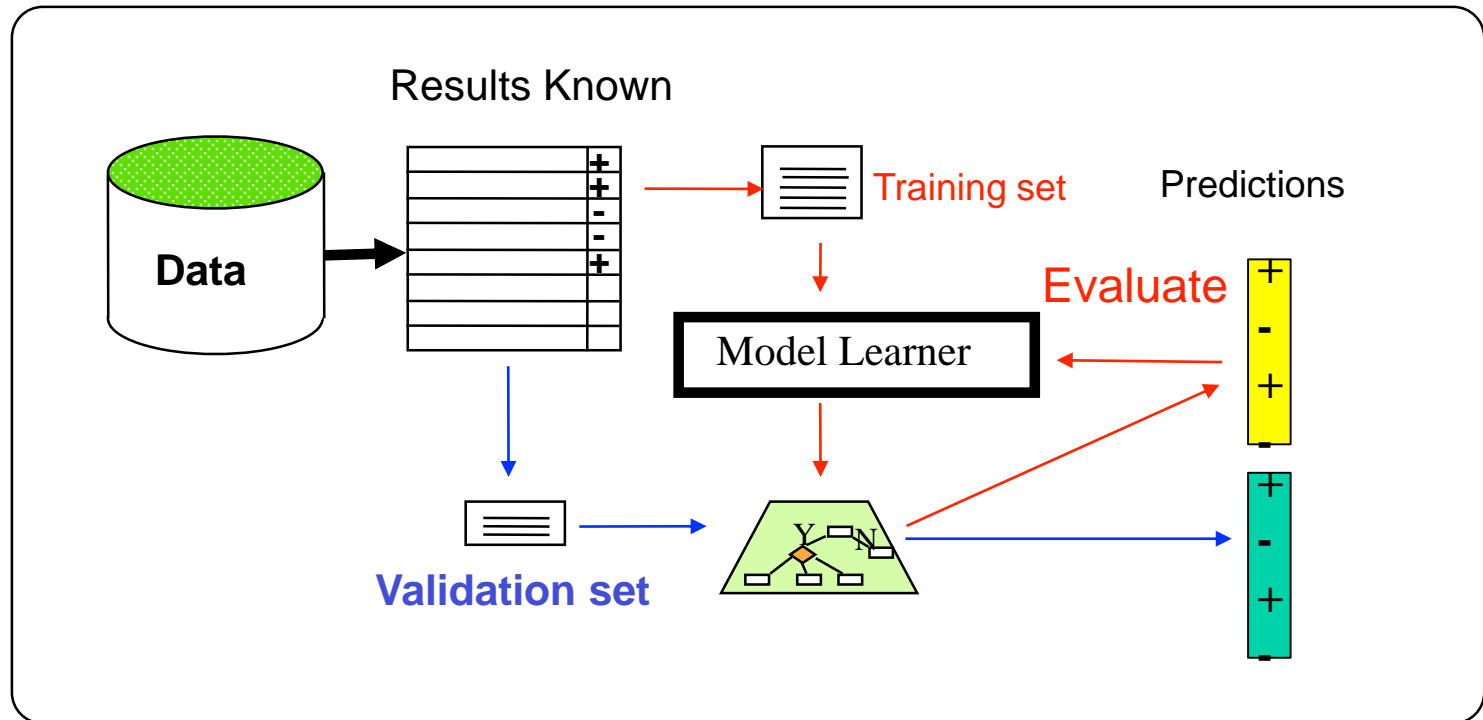
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

# Cross Validation

- **Cross Validation** is often used to counter overfitting.
- Partition the dataset into  $S$  groups, with  $(S-2)$  training sets, a validation set and a testing set.
  - The training set is used to determine the coefficients  $w$
  - The validation set is used to optimize the model complexity (hyperparameters, either  $M$  or  $\lambda$  in the previous example )
  - The testing set is used to evaluate the final selected mode
- The procedure is then repeated for all  $S$  possible choices, the performance scores from the  $S$  runs are then averaged.

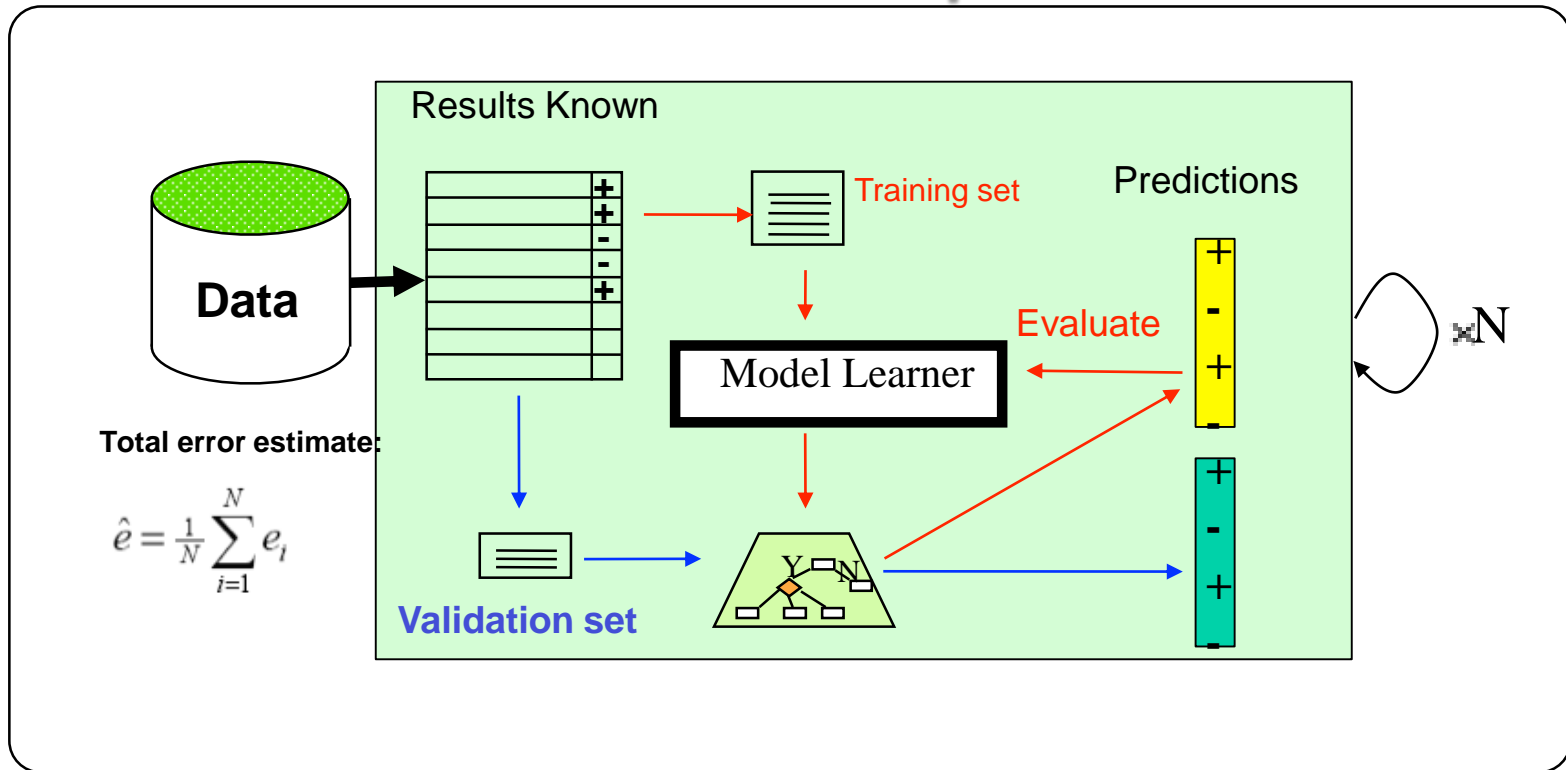


# Evaluation procedure



- For large datasets, a single split is usually sufficient.
- For smaller datasets, rely on cross validation

# Cross validation procedure



- Split the data into training, validation and test sets in a repeated fashion.
- Estimate the total error as the average of each fold error.

# Classification Measures - Error Rate

- Common performance measure for classification problems
  - Success: instance's class is predicted correctly (True Positives (**TP**) / Negatives (**TN**))
  - Error: instance's class is predicted incorrectly (False Positives (**FP**) / Negatives (**FN**))
  - False positives - **Type I error**. False Negative - **Type II error**
- Classification **error rate**: proportion of instances misclassified over the whole set of instances
$$\text{Error Rate} = (\text{FP} + \text{FN}) / (\text{P} + \text{N})$$
a.k.a, **accuracy**
$$\text{Accuracy} = 1 - \text{Error Rate}$$
- Classification Error Rate on the Training Set can be too optimistic!

# Unbalanced data

- Balanced set: (roughly) equal number of positive / negative examples:

Classifier	TP	TN	FP	FN	Recall Rate
A	25	25	25	25	50%
B	37	37	13	13	74%

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

# Unbalanced data

- Unbalanced set: unequal number of positive / negative examples

Classifier	TP	TN	FP	FN	Recall Rate
A	25	75	75	25	50%
B	0	150	0	50	0%




Classifier B cannot predict any positive examples!

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$


# Classification Measures

- Error Rate (Accuracy)
- Precision/Recall
- F-measure
- ROC curve
- Confusion matrix
- ...

# Recall/Precision



$$\text{RECALL: } \frac{A}{A+B} \times 100\%$$



$$\text{PRECISION: } \frac{A}{A+C} \times 100\%$$

More insight over a classifier's behavior

For the positive class:

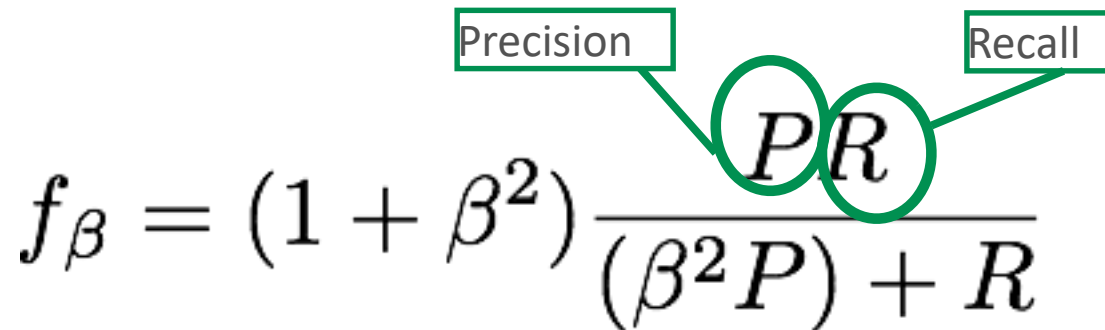
Classifier A: Recall = 50%, Precision = 25%

Classifier B: Recall = 0%, Precision = 0%

Classifier B is useless!

# F-measure

- Comparing different approaches is difficult when using multiple evaluation measures (e.g. Recall and Precision)
- F-measure combines recall and precision into a single measure:



The diagram shows the F-measure formula with annotations. The letters 'P' and 'R' in the numerator are each circled in green. A green box labeled 'Precision' has a line pointing to the 'P' circle, and another green box labeled 'Recall' has a line pointing to the 'R' circle.

$$f_{\beta} = (1 + \beta^2) \frac{PR}{(\beta^2 P) + R}$$


$\beta$  is a non-negative real values

- We often use f1 measure.



# ROC curves

- Receiver Operator Characteristic (ROC) curves plot TP vs FP rates



- Can be achieved by e.g. varying decision threshold of a classifier
- Area under the curve is often used as measure of goodness


# Confusion matrix

- A visualization tool used to present the results attained by a learner.
- Easy to see if the system is commonly mislabeling one class as another.

Predicted True	<b>A</b>	<b>B</b>	<b>C</b>
<b>A</b>	5	3	0
<b>B</b>	2	3	1
<b>C</b>	0	2	11

# Curse of Dimensionality

- As dimension  $D$  increases, the number of independent coefficients grows proportionally

$$y(x, \mathbf{w}) = w_0 + w_1 x^1 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$


$$y(x, \mathbf{w}) = w_0 + \sum_{i=1}^D w_i x_i + \sum_{i=1}^D \sum_{j=1}^D w_{ij} x_i x_j + \sum_{i=1}^D \sum_{j=1}^D \sum_{k=1}^D w_{ijk} x_i x_j x_k \dots$$

- The amount of data needed to support the result often grows exponentially with the dimensionality.

# Curse of Dimensionality

- Consider a sphere of radius  $r = 1$  in a space of  $D$  dimensions, what is the fraction of the volume of the sphere that lies between radius  $r = 1 - \epsilon$  and  $r = 1$ ?
- The volume of a sphere of radius  $r$  in  $D$  dimensions must scale as  $r^D$ , we have

$$V_3(r) = \frac{4}{3}\pi r^3$$


$$V_D(r) = K_D r^D$$

# Curse of Dimensionality

- Where the constant  $K_D$  depends only on  $D$ , thus the required volume fraction is given by

$$\frac{V_D(1) - V_D(1 - \epsilon)}{V_D(1)} = 1 - (1 - \epsilon)^D$$

- For large  $D$ , the fraction *tend to 1* even for small values of  $\epsilon$ .



# Curse of Dimensionality

- For large  $D$ , the fraction *tend to 1* even for small values of  $\epsilon$ .
- In space of high dimensionality, most of the volume of a sphere is concentrated in a *thin shell near the surface*.
- Distance functions losing their *usefulness* (for the nearest-neighbor criterion in feature-comparison algorithms, for example) in high dimensions.