



COMP3055


Machine Learning

Topic 12 – Multilayer Perceptrons

Zheng Lu
2023 Autumn


Limitations of Single Layer Perceptron

Only express linear decision surfaces



$$R = w_0 + \sum_{i=1}^n w_i x_i$$


$$o = \text{sign}(R) = \begin{cases} +1, & \text{if } R > 0 \\ -1, & \text{otherwise} \end{cases}$$



$$o = w_0 + \sum_{i=1}^n w_i x_i$$


Multilayer Perceptron (MLP)

- A more general network architecture: between the input and output layers there are hidden layers
- Hidden nodes do not directly receive inputs nor send outputs to the external environment
- Fully connected between layers




MLP Architecture

- Feedforward network: connections between the nodes do not form a cycle
- MLP usually interconnected in a feed-forward way
- The input layer does not count as a layer



3-layer feed-forward network

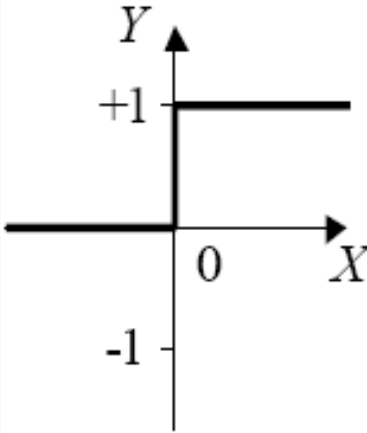
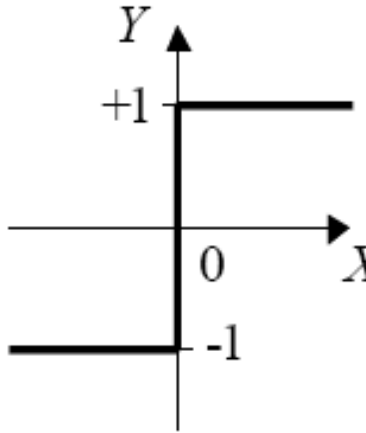
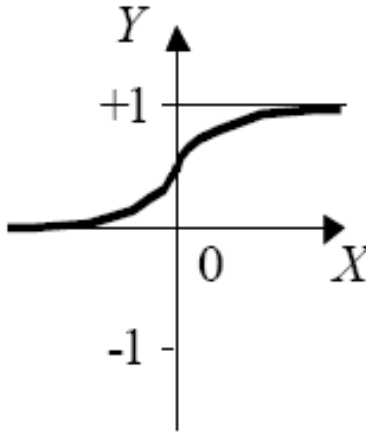
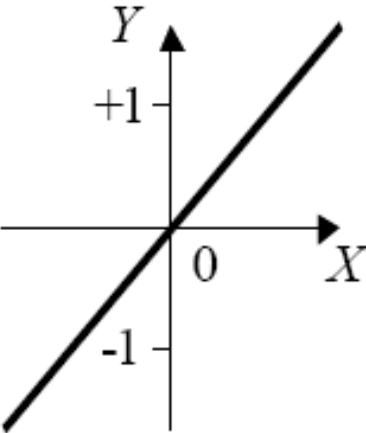


4-layer feed-forward network


Activation Function

- Activation functions are mathematical equations that determine the output of a neural network. The function is attached to each neuron in the network, and **determines whether it should be activated (“fired”) or not, based on whether each neuron’s input is relevant for the model’s prediction.** Activation functions also help normalize the output of each neuron to a range between 1 and 0 or between -1 and 1.
- The activation function can be considered as a mathematical “gate” in **between** the **input feeding the current neuron** and **its output going to the next layer.**

Activation Function

<i>Step function</i>	<i>Sign function</i>	<i>Sigmoid function</i>	<i>Linear function</i>
			
$Y^{step} = \begin{cases} 1, & \text{if } X \geq 0 \\ 0, & \text{if } X < 0 \end{cases}$	$Y^{sign} = \begin{cases} +1, & \text{if } X \geq 0 \\ -1, & \text{if } X < 0 \end{cases}$	$Y^{sigmoid} = \frac{1}{1 + e^{-X}}$	$Y^{linear} = X$

Sigmoid Unit




$\sigma(x)$ is the sigmoid function

$$\frac{1}{1 + e^{-x}}$$

Nice property: $\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$

We can derive gradient decent rules to train

Multilayer Perceptron



w_{ji} = weight associated with i th input to hidden unit j

w_{kj} = weight associated with j th input to output unit k

y_j = output of j th hidden unit

o_k = output of k th output unit

n = number of inputs

nH = number of hidden neurons

K = number of output neurons

$$y_j = \sigma\left(\sum_{i=0}^n x_i w_{ji}\right)$$

$$o_k = \sigma\left(\sum_{j=0}^{nH} y_j w_{kj}\right)$$

$$o_k = \sigma\left(\sum_{j=0}^{nH} \sigma\left(\sum_{i=0}^n x_i w_{ji}\right) w_{kj}\right)$$

Chain Rule

- In calculus, the chain rule is a formula to compute the derivative of a composite function.

Suppose that we have two functions $f(x)$ and $g(x)$ and they are both differentiable.

1. If we define $F(x) = (f \circ g)(x)$ then the derivative of $F(x)$ is,

$$F'(x) = f'(g(x)) g'(x)$$

2. If we have $y = f(u)$ and $u = g(x)$ then the derivative of y is,

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial f}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial z}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial q}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial y}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial y}$$

Chain rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial y}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2$, $y = 5$, $z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$



$$\frac{\partial f}{\partial x}$$

Chain Rule


$$f(x, y, z) = (x + y)z$$

e.g. $x = -2, y = 5, z = -4$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$




$$\frac{\partial f}{\partial x}$$

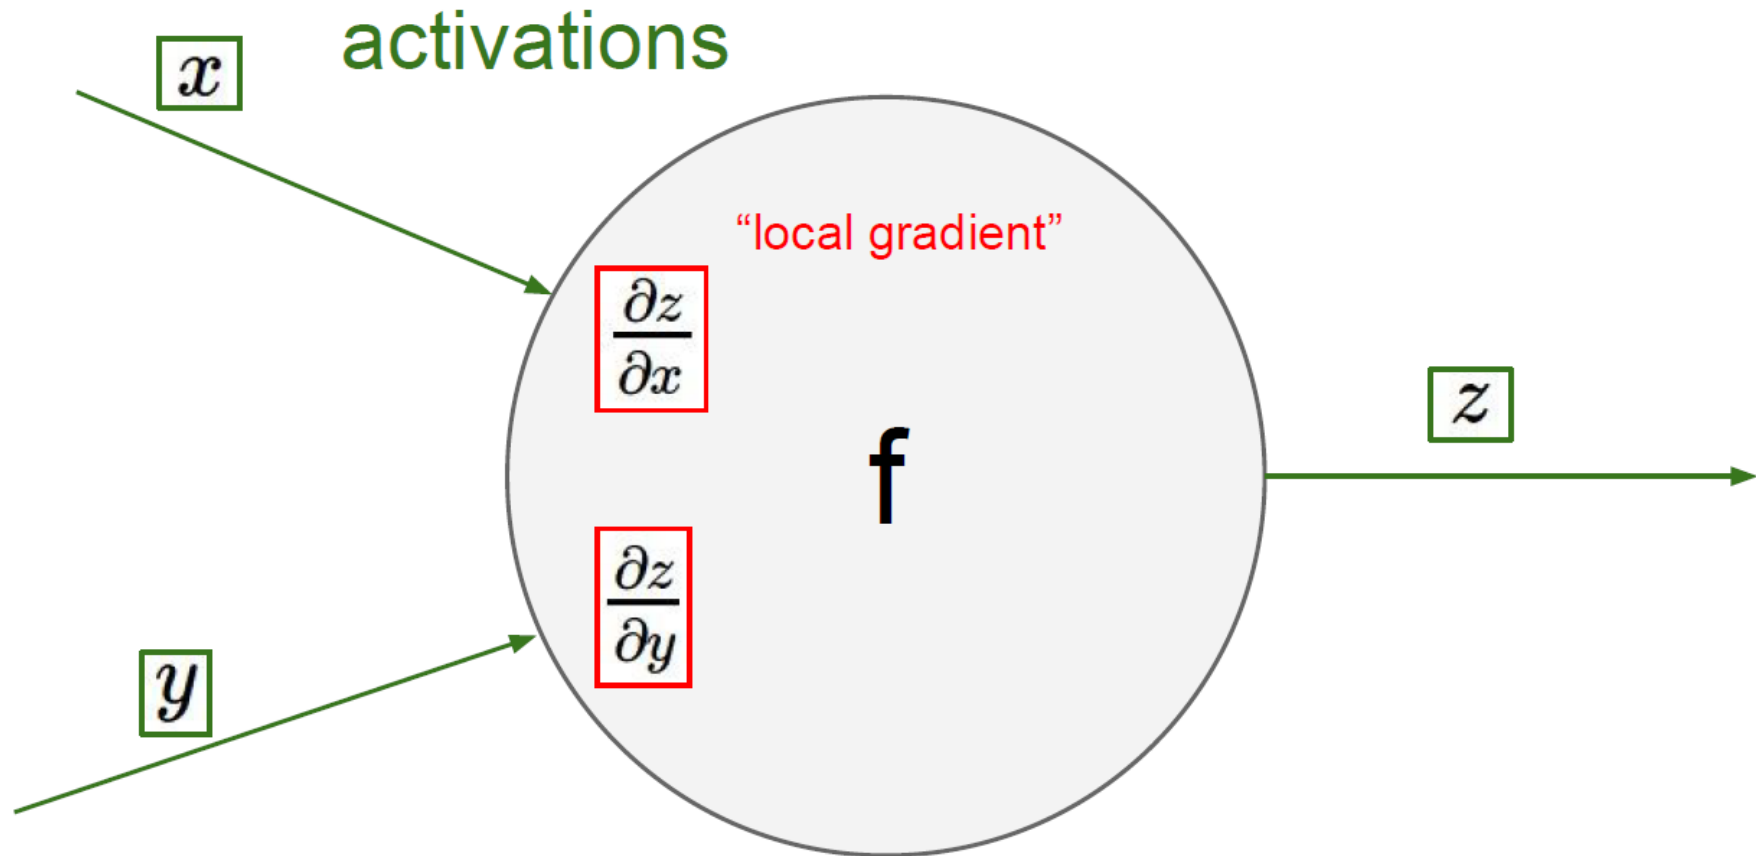
Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$


Chain Rule



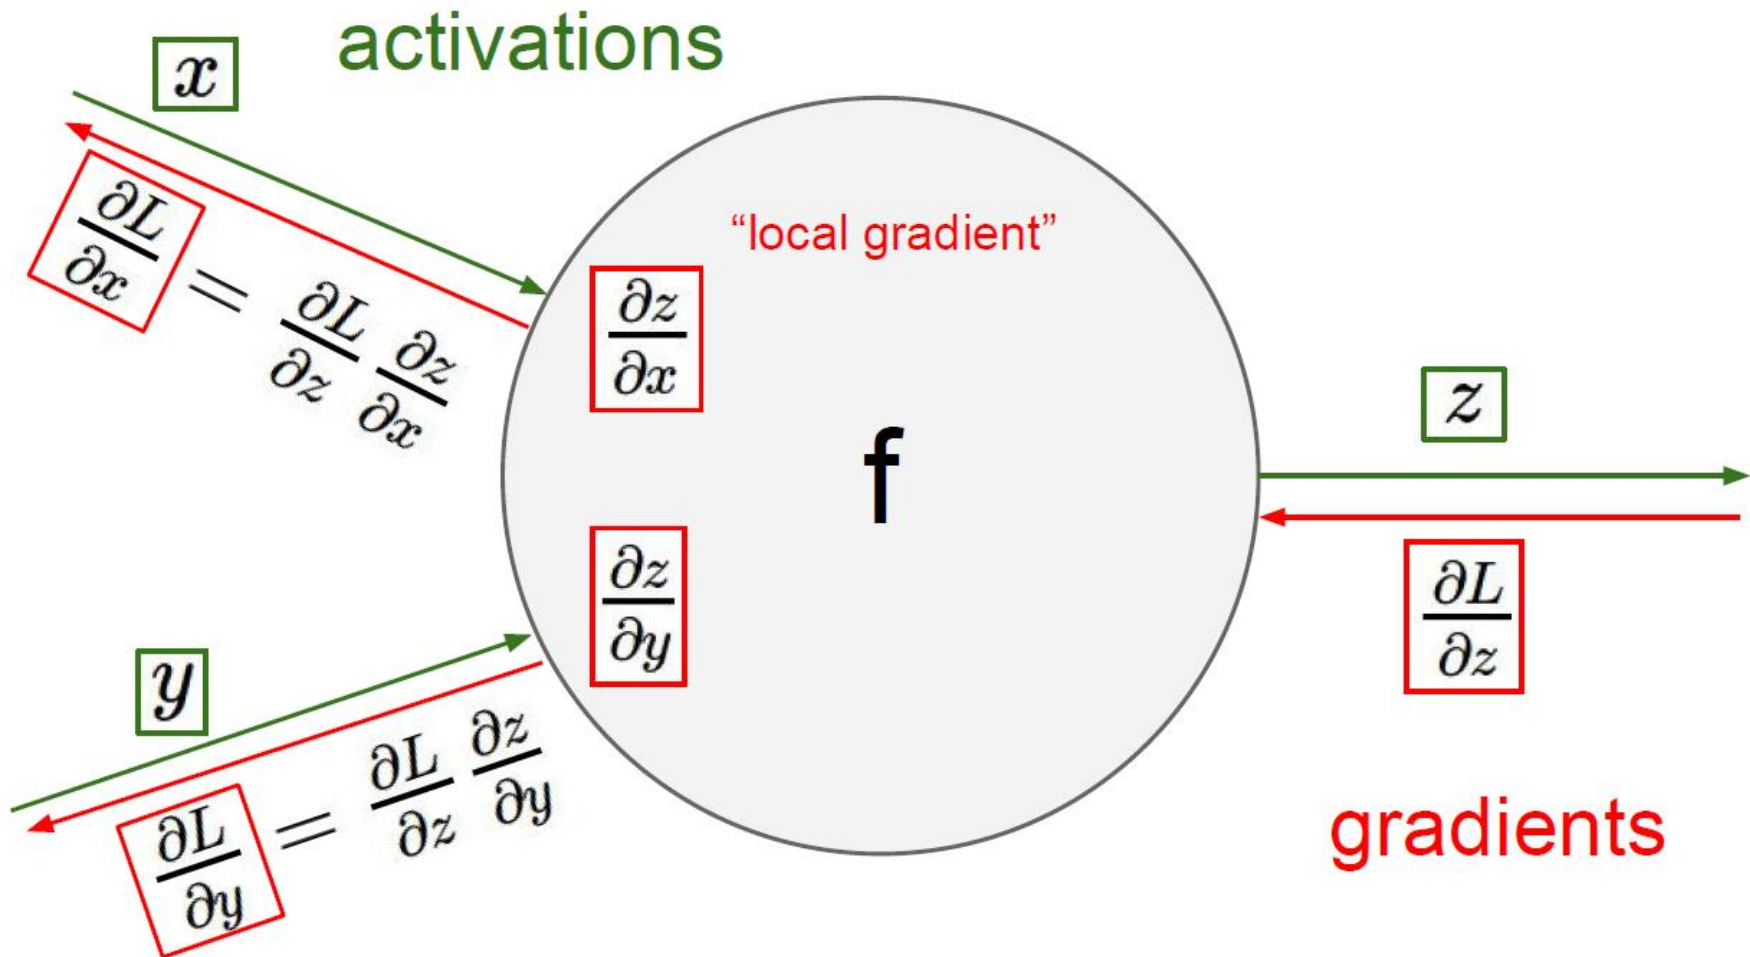
Chain Rule




Chain Rule



Chain Rule



Error Gradient for a Sigmoid Unit



$$E(W) \equiv \frac{1}{2} \sum_{k=1}^K (d(k) - o(k))^2$$

Error Gradient for a Sigmoid Unit

$$E(W) \equiv \frac{1}{2} \sum_{k=1}^K (d(k) - o(k))^2$$

$$\begin{aligned} \frac{\partial E}{\partial w_i} &= \frac{\partial E}{\partial w_i} \left(\frac{1}{2} \sum_{k=1}^K (d(k) - o(k))^2 \right) \\ &= \frac{1}{2} \sum_{k=1}^K \left(\frac{\partial E}{\partial w_i} (d(k) - o(k))^2 \right) \\ &= \frac{1}{2} \sum_{k=1}^K \left(2(d(k) - o(k)) \frac{\partial E}{\partial w_i} (d(k) - o(k)) \right) \\ &= \sum_{k=1}^K \left((d(k) - o(k)) \frac{\partial E}{\partial w_i} (-(o(k))) \right) \\ &= - \sum_{k=1}^K \left((d(k) - o(k)) \frac{\partial o(k)}{\partial net(k)} \frac{\partial net(k)}{\partial w_i} \right) \end{aligned}$$

Error Gradient for a Sigmoid Unit

$$\frac{\partial o(k)}{\partial net(k)} = \frac{\partial \sigma(net(k))}{\partial net(k)} = \sigma(net(k)) (1 - \sigma(net(k))) = o(k)(1 - o(k))$$

$$\frac{\partial net(k)}{\partial w_i} = \frac{\partial (WX(k))}{\partial w_i} = x_i(k)$$


$$\begin{aligned} \frac{\partial E}{\partial w_i} &= - \sum_{k=1}^K \left((d(k) - o(k)) \frac{\partial o(k)}{\partial net(k)} \frac{\partial net(k)}{\partial w_i} \right) \\ &= - \sum_{k=1}^K \left((d(k) - o(k)) o(k) (1 - o(k)) x_i(k) \right) \end{aligned}$$

Back-propagation: Initial Steps

- Training Set: A set of input vectors x_i , $i=1...D$ with the corresponding targets t_i .
- η : learning rate, controls the change rate of the weights.
- Begin with random weights.


Back-propagation: Output Neurons

$$E(W) \equiv \frac{1}{2} \sum_{k=1}^K (d(k) - o(k))^2$$




- E depends on the weights because $o_k = \sigma\left(\sum_{j=1}^{nH} y_j w_{kj}\right)$
- For simplicity we assume the error of one training example

Back-propagation: Output Neurons




- $\frac{\partial E_k}{\partial w_{kj}} = \frac{\partial E_k}{\partial o_k} \frac{\partial o_k}{\partial net_k} \frac{\partial net_k}{\partial w_{kj}} = \frac{\partial E_k}{\partial o_k} \frac{\partial \sigma(net_k)}{\partial net_k} y_j$
- We define $\delta_k = \frac{\partial E_k}{\partial o_k} \frac{\partial \sigma(net_k)}{\partial net_k}$
- Update: $\Delta w_{kj} = -\eta \frac{\partial E_k}{\partial w_{kj}} = -\eta \delta_k y_j$

Back-propagation: Hidden Neurons




- $\frac{\partial E}{\partial w_{ji}} = \frac{\partial E}{\partial y_j} \frac{\partial y_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}} = \frac{\partial E}{\partial y_j} \frac{\partial \sigma(net_j)}{\partial net_j} x_i$
- $\frac{\partial E}{\partial y_j} = \sum_{k=1}^K \frac{\partial E_k}{\partial y_j} = \sum_{k=1}^K \frac{\partial E_k}{\partial o_k} \frac{\partial o_k}{\partial y_j} = \sum_{k=1}^K \frac{\partial E_k}{\partial o_k} \frac{\partial \sigma(net_k)}{\partial net_k} \frac{\partial net_k}{\partial y_j}$

Back-propagation: Hidden Neurons




- $$\frac{\partial E}{\partial y_j} = \sum_{k=1}^K \frac{\partial E_k}{\partial o_k} \frac{\partial o_k}{\partial net_k} \frac{\partial net_k}{\partial y_j} = \sum_{k=1}^K \delta_k w_{kj}$$
- $$\frac{\partial E}{\partial w_{ji}} = \frac{\partial E}{\partial y_j} \frac{\partial y_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}} = \sum_{k=1}^K (\delta_k w_{kj}) \frac{\partial \sigma(net_j)}{\partial net_j} x_i$$

Back-propagation: Hidden Neurons



- $\frac{\partial E}{\partial w_{ji}} = \sum_{k=1}^K (\delta_k w_{kj}) \frac{\partial \sigma(net_j)}{\partial net_j} x_i$
- We define $\delta_j = \sum_{k=1}^K (\delta_k w_{kj}) \frac{\partial \sigma(net_j)}{\partial net_j}$

Back-propagation: Hidden Neurons



- $\frac{\partial E}{\partial w_{ji}} = \delta_j x_i$
- Update: $\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}} = -\eta \delta_j x_i$

Back-propagation Summary

1. Initialise weights randomly
2. For each input training example x compute the outputs (**forward pass**)
3. Compute the output neurons errors and then compute the update rule for output layer weights (**backward pass**)

$$\Delta w_{kj} = -\eta \frac{\partial E}{\partial w_{kj}} = -\eta \delta_k y_j \text{ where } \delta_k = \frac{\partial E}{\partial o_k} \frac{\partial \sigma(net_k)}{\partial net_k}$$

4. Compute hidden neurons errors and then compute the update rule for hidden layer weights (**backward pass**)

$$\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}} = -\eta \delta_j x_i \text{ where } \delta_j = \sum_{k=1}^K (\delta_k w_{kj}) \frac{\partial \sigma(net_j)}{\partial net_j}$$

Back-propagation Summary

5. Compute the sum of all Δw , once all training examples have been presented to the network
 6. Update weights $w_i \leftarrow w_i + \Delta w_i$
 7. Repeat steps 2-6 until the stopping criterion is met
- The algorithm will converge to a weight vector with minimum error, given that the learning rate is sufficiently small

Back-propagation Summary

- Gradient descent over entire network weight vector.
- Will find a local, not necessarily a global error minimum.
- In practice, it often works well (can run multiple times).
- Minimizes error over all training samples.
 - Will it generalize to subsequent examples? i.e., will the trained network perform well on data outside the training sample.
- Training can take thousands of iterations.
- After training, use the network is fast.



Generalization, Overfitting and Stopping Criterion

What is the appropriate condition for stopping weight update?

Continue until the error E falls below some predefined value?


- Not a very good idea.
- Back-propagation is susceptible to overfitting the training example at the cost of decreasing generalization accuracy over other unseen examples.

Generalization, Overfitting and Stopping Criterion




- Stop training when the validation set has the lowest error
- Error might decrease in the training set but increase in the 'validation' set (overfitting!)
- Early stopping: one way to avoid overfitting

Dropout



- **Dropout:** Randomly remove some nodes in the network (along with incoming and outgoing edges)
- Notes:
 - Usually $p \geq 0.5$ (p is probability of keeping node)
 - Input layers p should be much higher (and use noise instead of dropout)
 - Most deep learning frameworks come with a dropout layer

Weight Decay



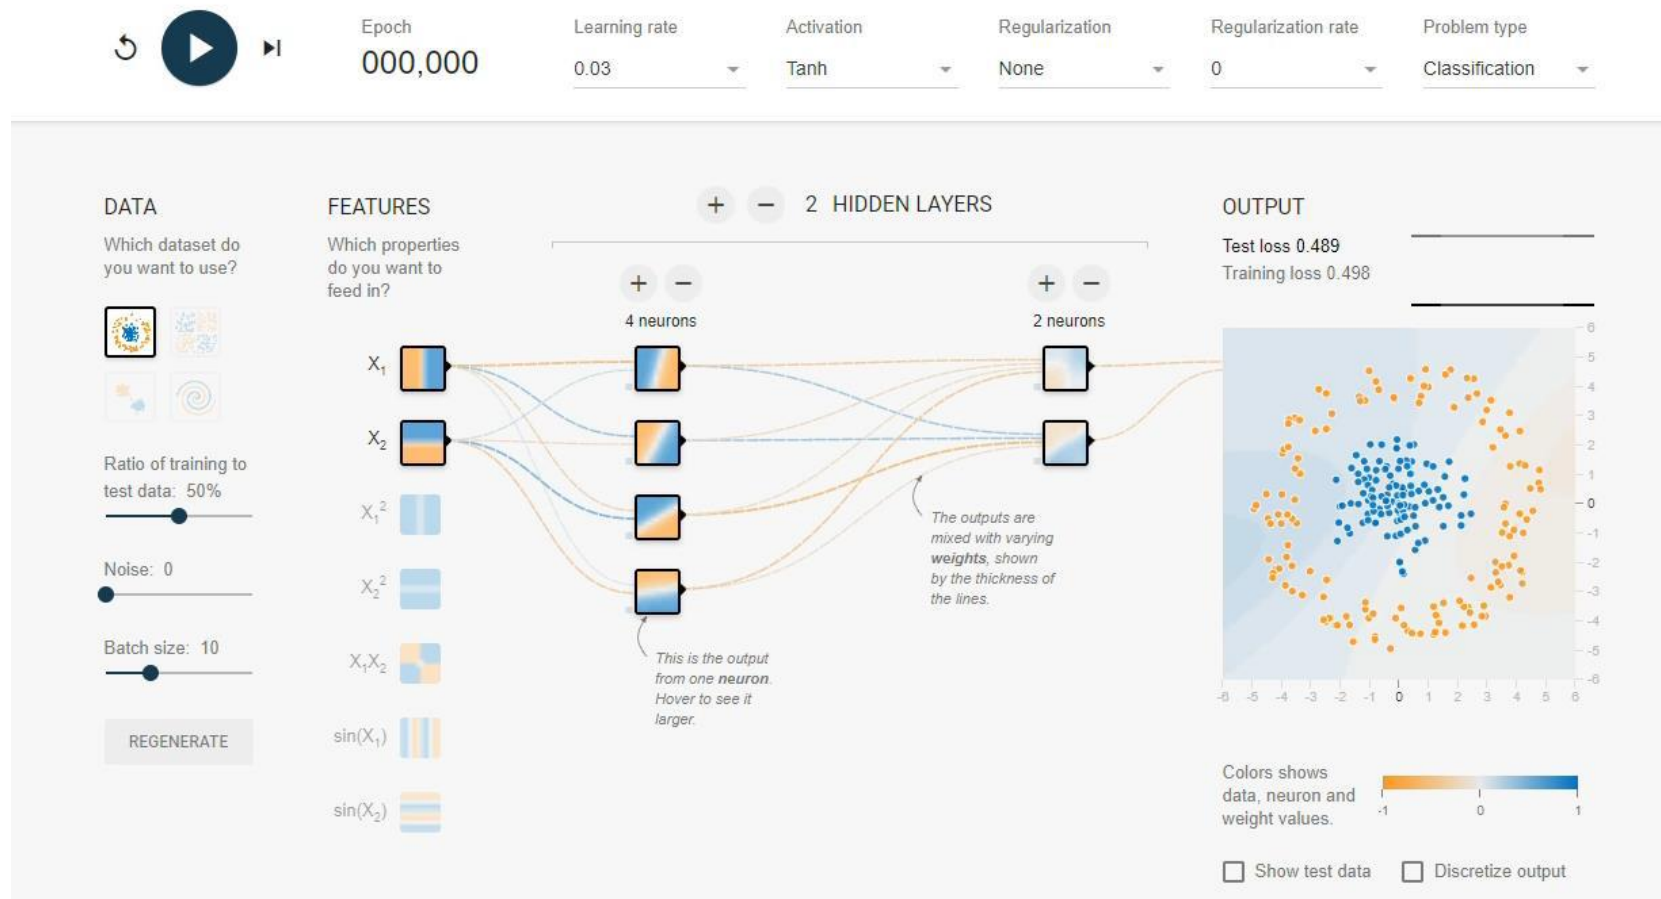
- **L2 Penalty:** Penalize squared weights. Result:
 - Keeps weight small unless error derivative is very large.
 - Prevent from fitting sampling error.
 - Smoother model (output changes slower as the input change).
 - If network has two similar inputs, it prefers to put half the weight on each rather than all the weight on one.
- **L1 Penalty:** Penalize absolute weights. Result:
 - Allow for a few weights to remain large.

Normalization

- Network Input Normalization
 - *Example:* Pixel to $[0, 1]$ or $[-1, 1]$ or according to mean and std.
- Batch Normalization (BatchNorm, BN)
 - Normalize hidden layer inputs to mini-batch mean & variance
 - Reduces impact of earlier layers on later layers
- Batch Renormalization (BatchRenorm, BR)
 - Fixes difference b/w training and inference by keeping a moving average asymptotically approaching a global normalization.

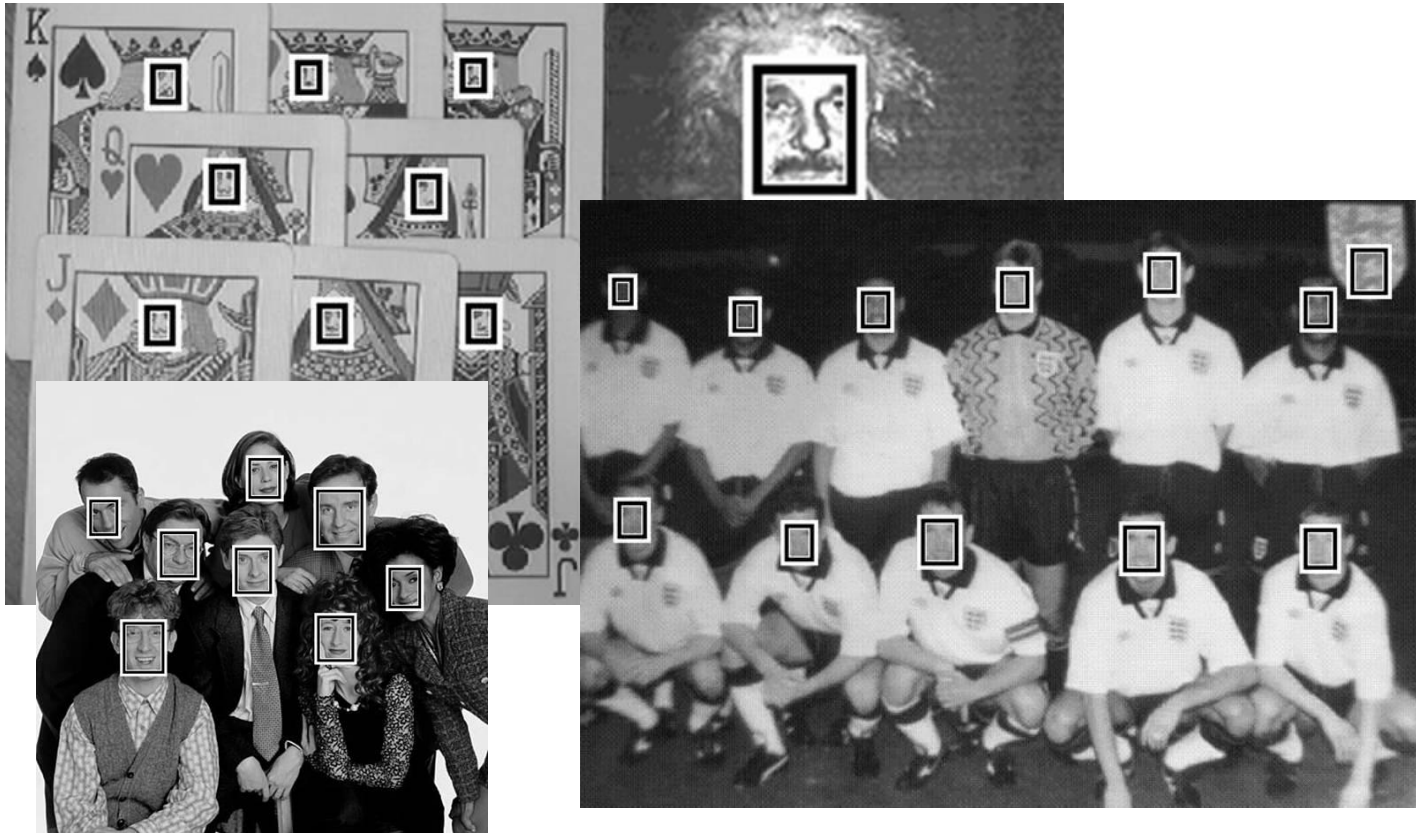
Neural Network Playground

<http://playground.tensorflow.org>



Application Example

Neural Network-based Face Detection



https://ri.cmu.edu/pub_files/pub1/rowley_henry_1996_3/rowley_henry_1996_3.pdf

Application Example


Neural Network-based Face Detection



Application Example

Neural Network-based Face Detection

- It takes 20 x 20 pixel window, feeds it into a NN, which outputs a value ranging from -1 to $+1$ signifying the presence or absence of a face in the region.
- The window is applied at every location of the image.
- To detect faces larger than 20 x 20 pixel, the image is repeatedly reduced in size.



Application Example

Neural Network-based Face Detection

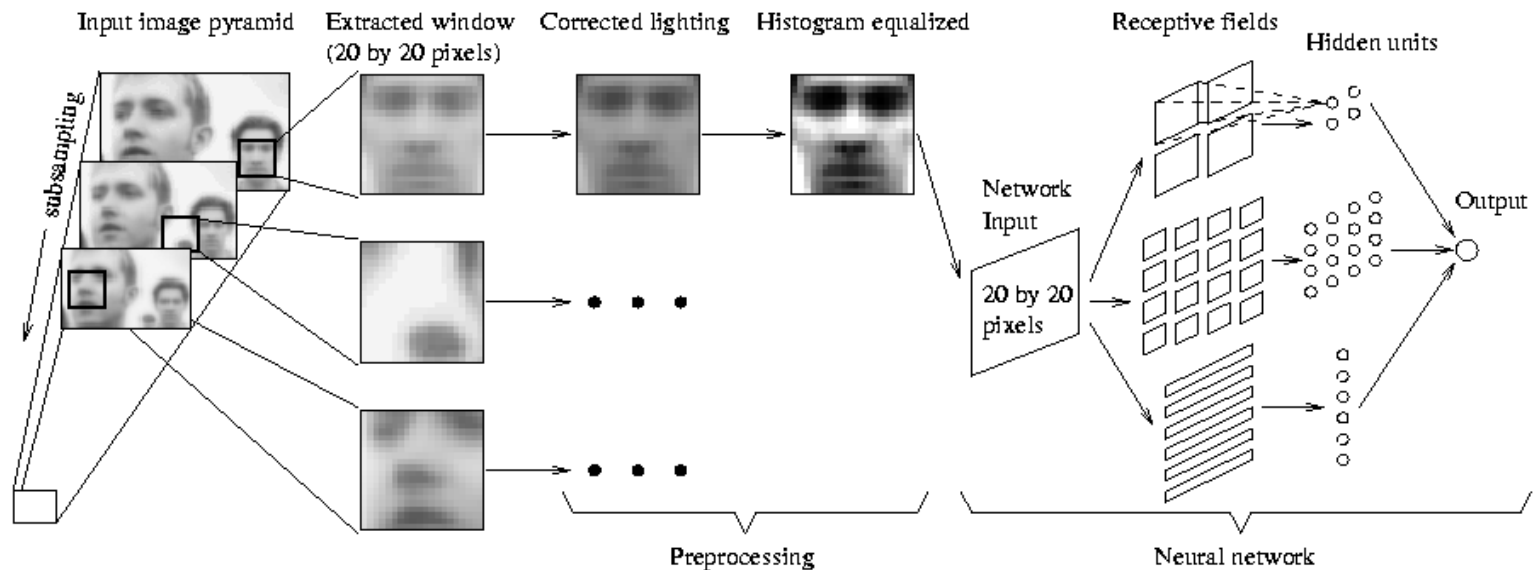


Figure 1: The basic algorithm used for face detection.

Application Example

Neural Network-based Face Detection




Figure 1: The basic algorithm used for face detection.

- 4 look at 10 x 10 subregions
- 16 look at 5x5 subregions
- 6 look at 20x5 horizontal stripes of pixels

Application Example

Neural Network-based Face Detection




Figure 1: The basic algorithm used for face detection.

- Training samples
- 1050 initial face images. More face example are generated from this set by rotation and scaling. Desired output +1
- Non-face training samples: Use a bootstrapping technique to collect 8000 non-face training samples from 146,212,178 subimage regions! Desired output -1

Application Example

Neural Network-based Face Detection

- Training samples: Non-face training samples





Figure 1: The basic algorithm used for face detection.



Application Example

Neural Network-based Face Detection

- Post-processing and face detection





Figure 1: The basic algorithm used for face detection.



Application Examples

Neural Network-based Face Detection

- Results and Issues

- 77.% ~ 90.3% detection rate (130 test images)
- Process 320x240 image in 2 – 4 seconds on a 200MHz R4400 SGI Indigo 2




Figure 1: The basic algorithm used for face detection.

Further Reading

Chapter 4, T. M. Mitchell, Machine Learning,
McGraw-Hill International Edition, 1997