

COMP3052.SEC Computer Security

Session 07: Reference Monitors

```
ide1: BM-DMA at 0xc008-0xc00f, BIOS settings: hdc:pio, hdd:pio
ne2k-pci.c:v1.03 9/22/2003 D. Becker/P. Gortmaker
  http://www.scyld.com/network/ne2k-pci.html
hda: QEMU HARDDISK, ATA DISK drive
ide0 at 0x1f0-0x1f7,0x3f6 on irq 14
hdc: QEMU CD-ROM, ATAPI CD/DVD-ROM drive
ide1 at 0x170-0x177,0x376 on irq 15
ACPI: PCI Interrupt Link [LNKC] enabled at IRQ 10
ACPI: PCI Interrupt 0000:00:03.0[A] -> Link [LNKC] -> GSI 10 (level, low) -> IRQ
  10
eth0: RealTek RTL-8029 found at 0xc100, IRQ 10, 52:54:00:12:34:56.
hda: max request size: 512KiB
hda: 180224 sectors (92 MB) w/256KiB Cache, CHS=178/255/63, (U)DMA
hda: set_multimode: status=0x41 { DriveReady Error }
hda: set_multimode: error=0x04 { DriveStatusError }
ide: failed opcode was: 0xef
hda: cache flushes supported
  hda: hda1
hdc: ATAPI 4X CD-ROM drive, 512kB Cache, (U)DMA
Uniform CD-ROM driver Revision: 3.20
Done.
Begin: Mounting root file system... ..
/init: /init: 151: Syntax error: 0xforce=panic
Kernel panic - not syncing: Attempted to kill init!
```

Acknowledgements

- Some of the materials we use this semester may come directly from previous teachers of this module, and other sources ...
- Thank you to (amongst others):
 - Michel Valstar, Milena Radenkovic, Mike Pound, Dave Towey, ...


This Session

- Reference Monitors
- Operating System Integrity
- Privilege Elevation
- Memory Protection
- Page Tables

Concepts

- The Reference Monitor
 - An abstract concept
- Security Kernel
 - The implementation of a reference monitor
- Trusted Computing Base (TCB)
 - Kernel + other protection measures

Reference Monitor



“An access control concept that refers to an abstract machine that mediates all access to objects by subjects”

- Must be tamper proof / resistant
- **Must always be invoked** when access to an object is required
- Must be small enough to be verifiable / subject to analysis to ensure correctness

Security Kernel

“The hardware, firmware and software elements of a TCB that implement the reference monitor”

- Mediates all access
- Must be protected from modification
- Must be verifiably correct
- Usually in the bottom layers of a system


Trusted Computing Base

“The totality of protection mechanisms within a computer system responsible for enforcing a security policy”

- One or more components
- Enforce a unified security policy over a product or system
- Correct enforcement depends on components within as well as input from administrators

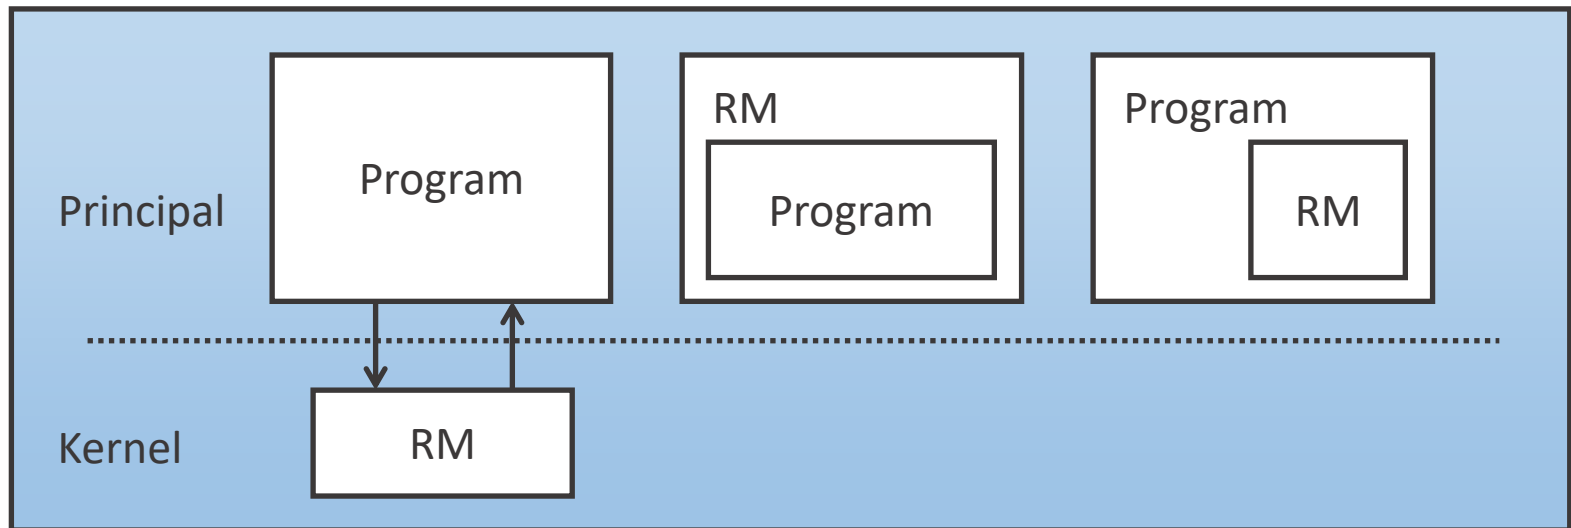
Placement

- Can be placed anywhere within a system
 - Hardware – Dedicated registers for defining privileges
 - Operating system kernel – E.g. Virtual Machine Hypervisor
 - Operating system – Windows security reference monitor
 - Services Layer – JVM, .NET
 - Application Layer – Firewalls



Placement

- Reference monitors could be placed in a variety of locations relative to the program being run



Lower Is Better

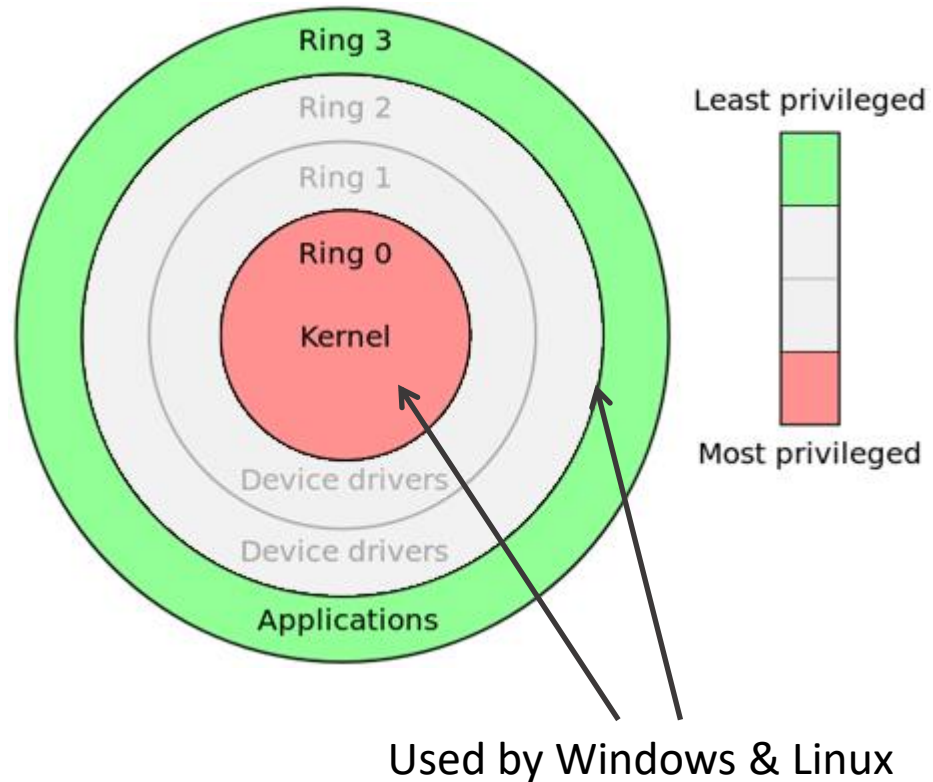
- Using a reference monitor or other security features at a lower level means:
 - We can assure a higher degree of security
 - Usually simple structures to implement
 - Reduced performance overheads
 - Fewer layer below attack possibilities
- However:
 - Access control decisions are far remote from applications

OS Integrity


- The operating system:
 - Arbitrates access requests
 - Is itself a resource that must be accessed
- This is a conflict, we want to use the OS but not mess with it
“Users must not be able to modify the operating system”
- Modes of operation
 - Defines which actions are permitted in which mode, e.g. system calls, machine instructions, I/O
- Controlled Invocation
 - Allows us to execute privileged instructions safely, before returning to user code

Modes of Operation

- Distinguish between computations done on behalf of:
 - The OS
 - The user
- A status flag within the CPU allows the OS to operate in different modes



Controlled Invocation

- Many functions are held at kernel level, but are quite reasonably called from within user level code
 - Network and File IO
 - Memory allocation
 - Halting the CPU (at shutdown only!)
 - We need a mechanism to transfer between kernel mode (ring 0) and user mode (ring 3)
- 
- Protection Fault

The Key Point

We don't actually perform privileged operations: we ask the OS to perform them for us

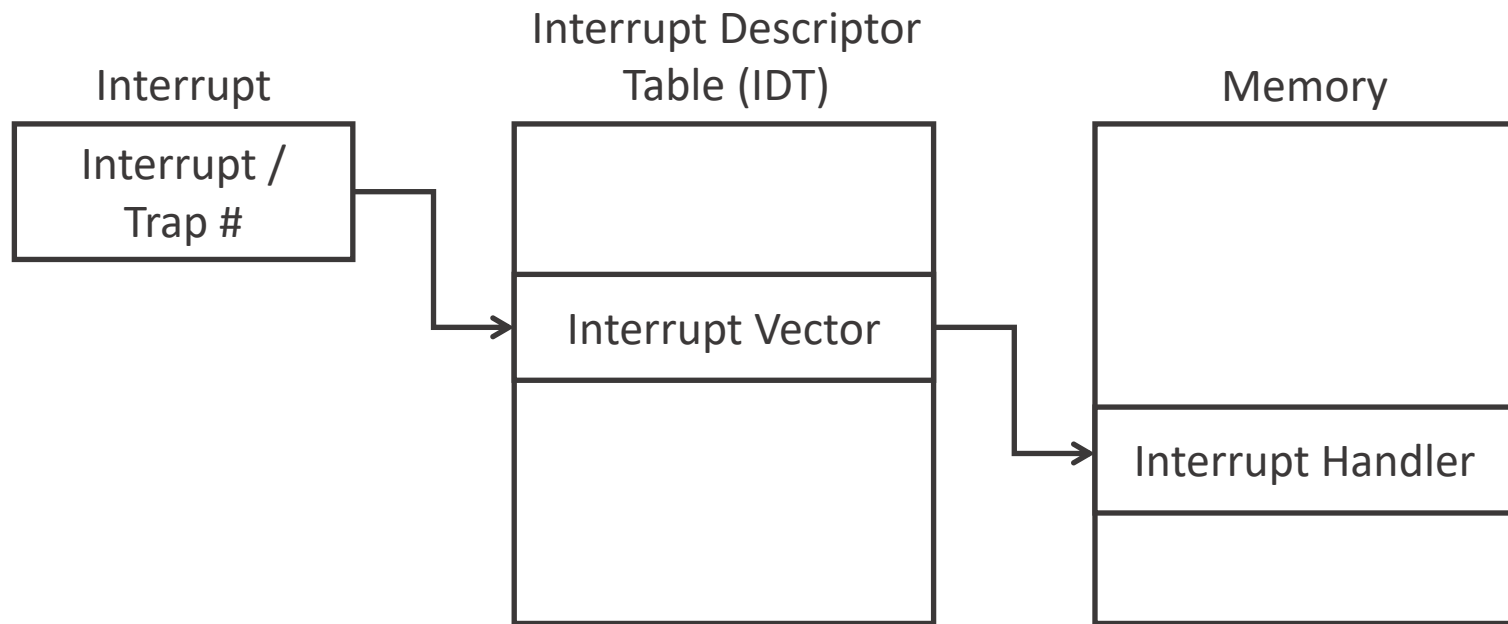
The OS can refuse to do it!

Controlled Invocation: Interrupts

- Exceptions / Interrupts / Traps
 - Called various things, for now we'll just use "Interrupt"
 - In many ways is the hardware equivalent to a software exception
- Handled by an interrupt handler which resolves the issue and returns to the original code

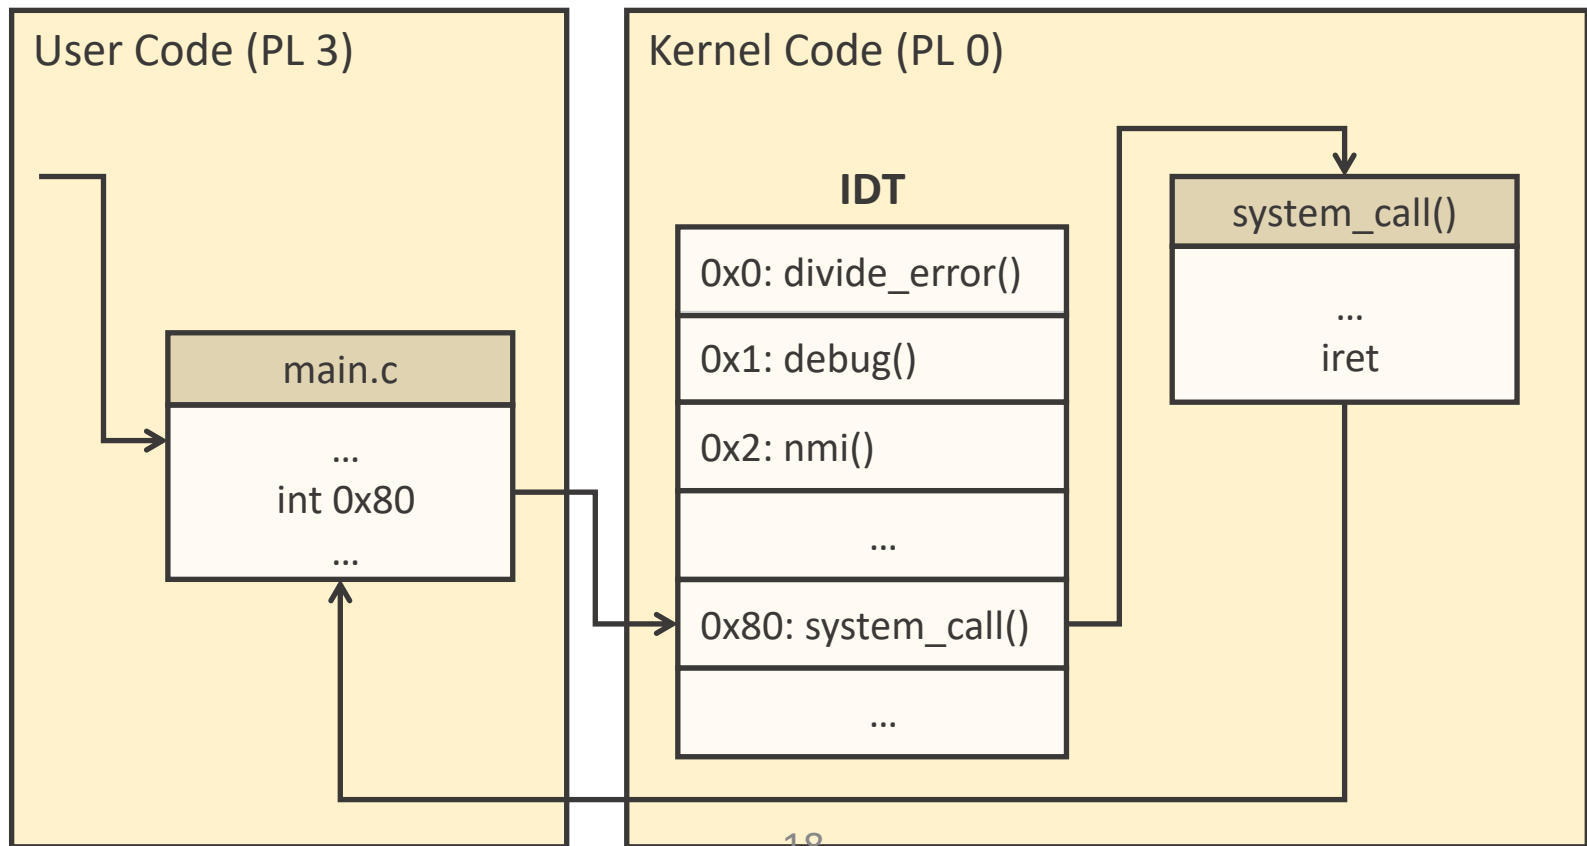
Processing an Interrupt

- Given an interrupt, the CPU will switch execution to the location given in an interrupt descriptor table (IDT)



Privilege Elevation in x86-Linux

- Linux initialises its IDT to handle syscalls at vector 0x80



Interrupt Example

```
#include <unistd.h>

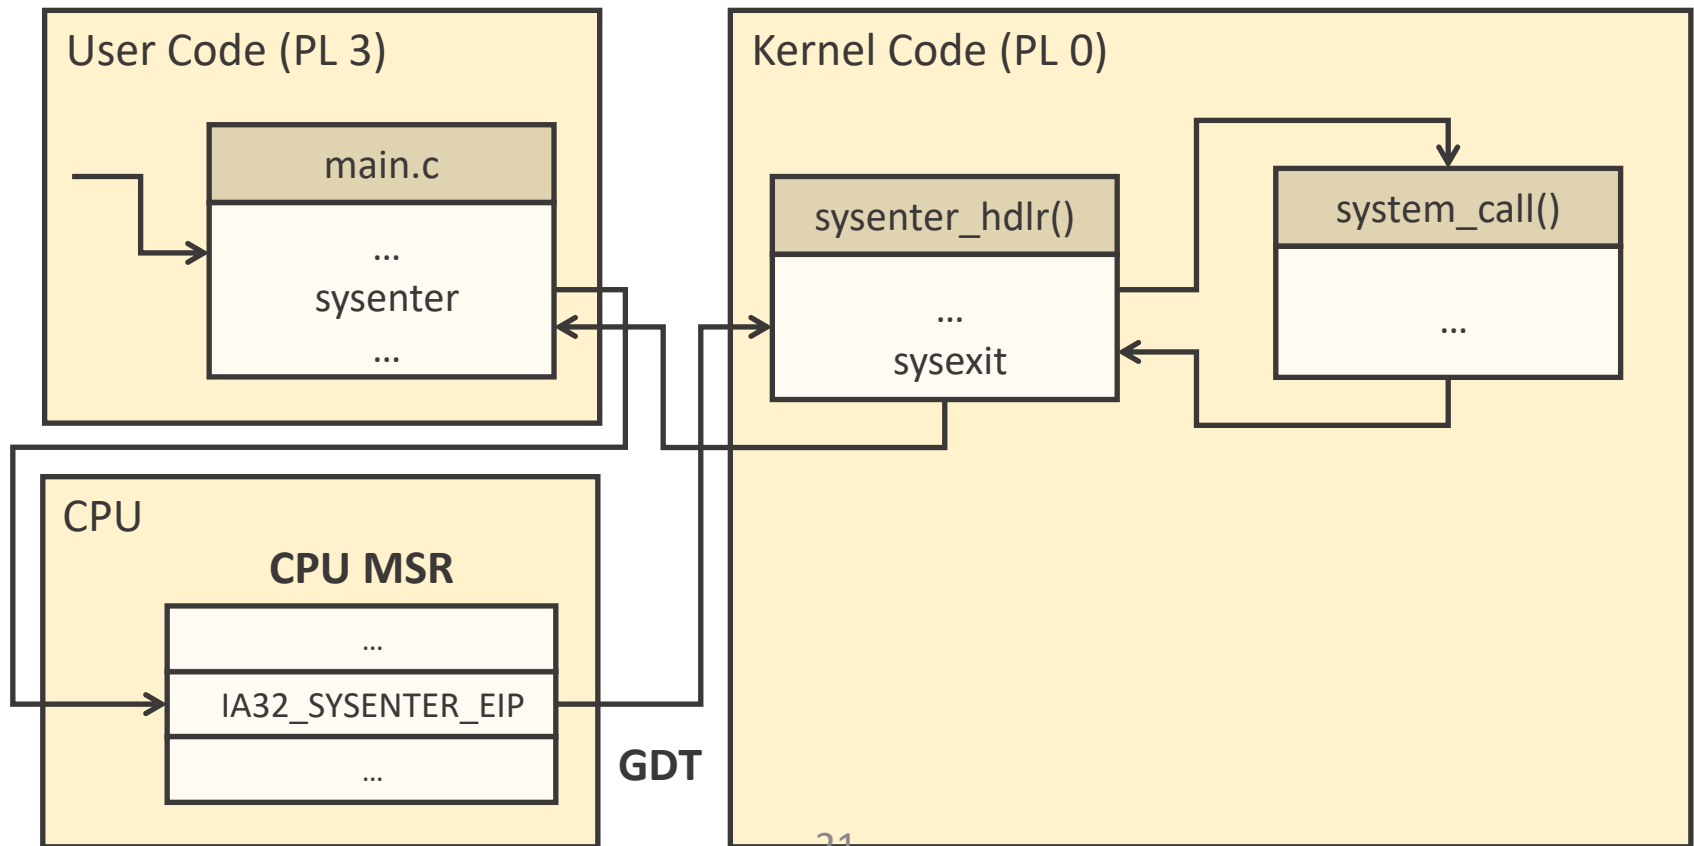
int main(int argc, char *argv[])
{
    write(1, "Hello!", 6);
    _exit(0);
}
```

```
mov    $4, %eax
mov    $1, %ebx
mov    $msg, %ecx
mov    $6, %edx
int    $0x80

mov    $1, %eax
mov    $0, %ebx
int    $0x80
```

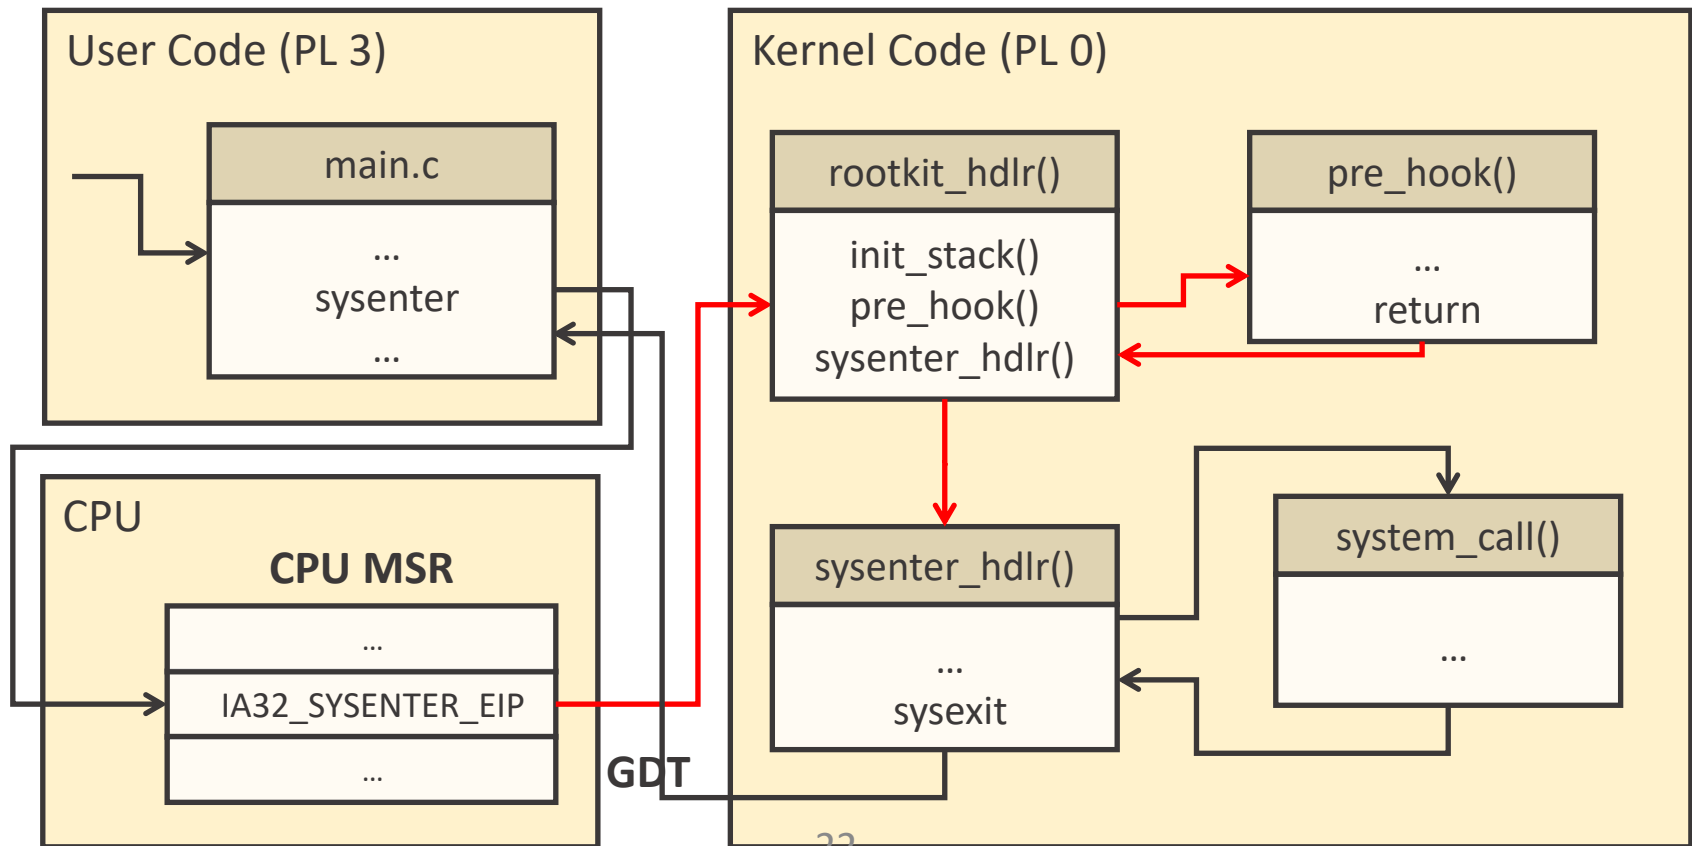
Modern Kernels

- Intel introduced the `sysenter` and `sysexit` operations with the Pentium II – much less overhead



Patching the Kernel

- If you can run custom PL 0 code (compromised driver?), you can insert your own handler – **Rootkit**



Processes and Threads

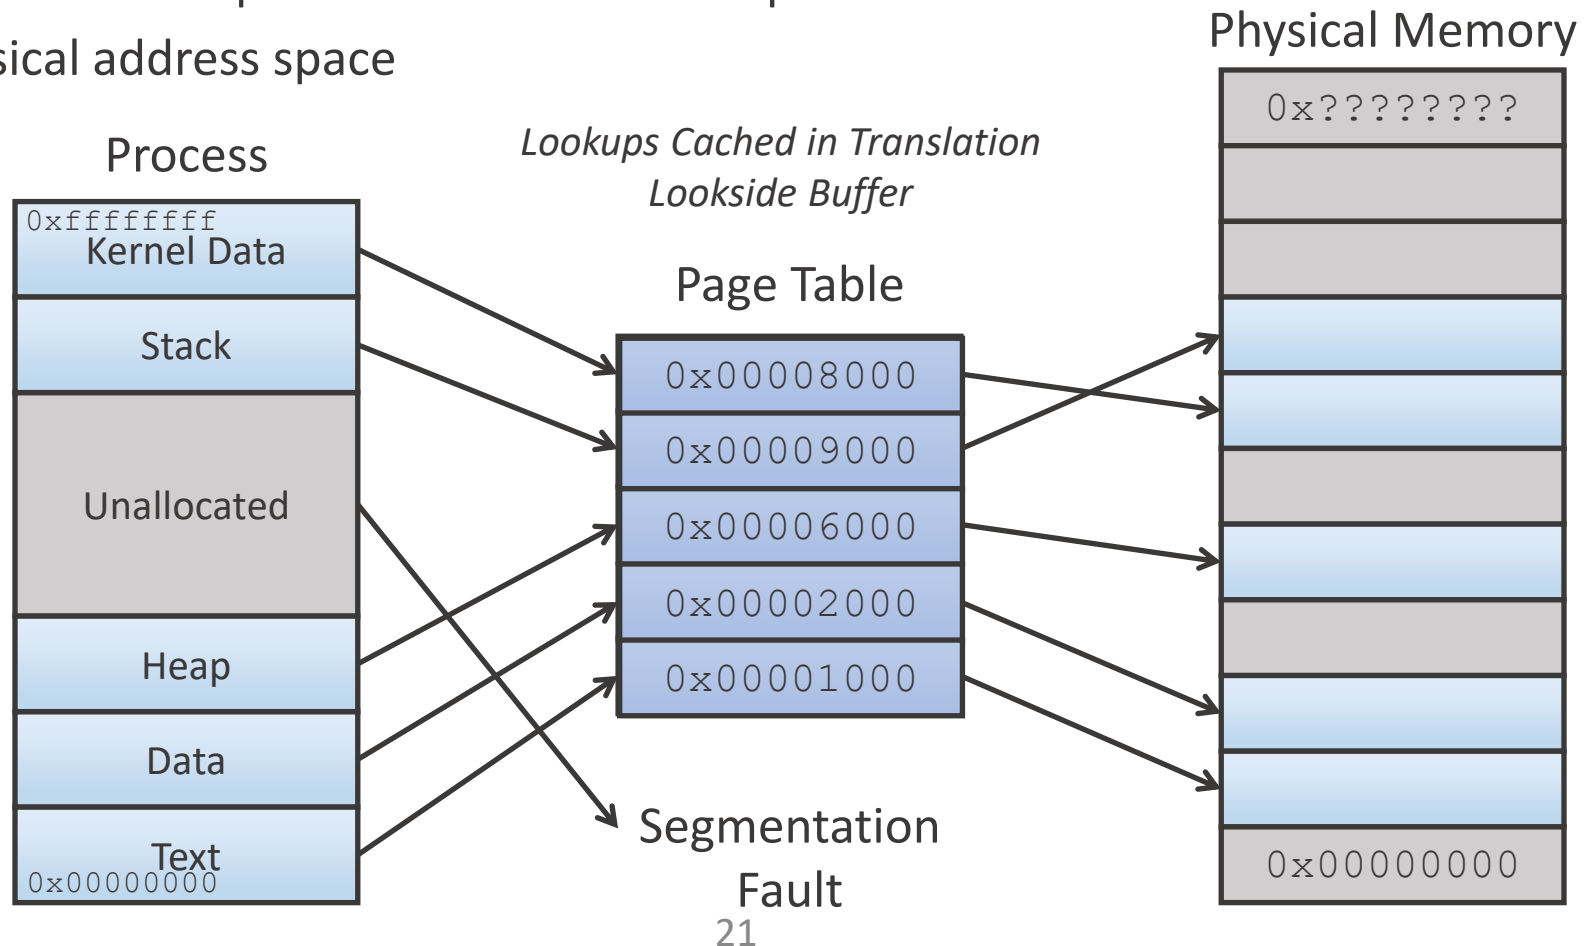
- A process is a program being executed
- Important unit of control:
 - Exists in its own address space
 - Communicates with other processes via the OS
 - Separation for security
- A Thread is a strand of execution within a process
 - Share a common address space

Memory Protection

- Segmentation – divides data into logical units
 - Good for security
 - Challenging memory management
 - Not used much in modern OSs
- Paging – divides memory into pages of equal size
 - Efficient memory management
 - Less good for access control
 - Extremely common in modern OSs

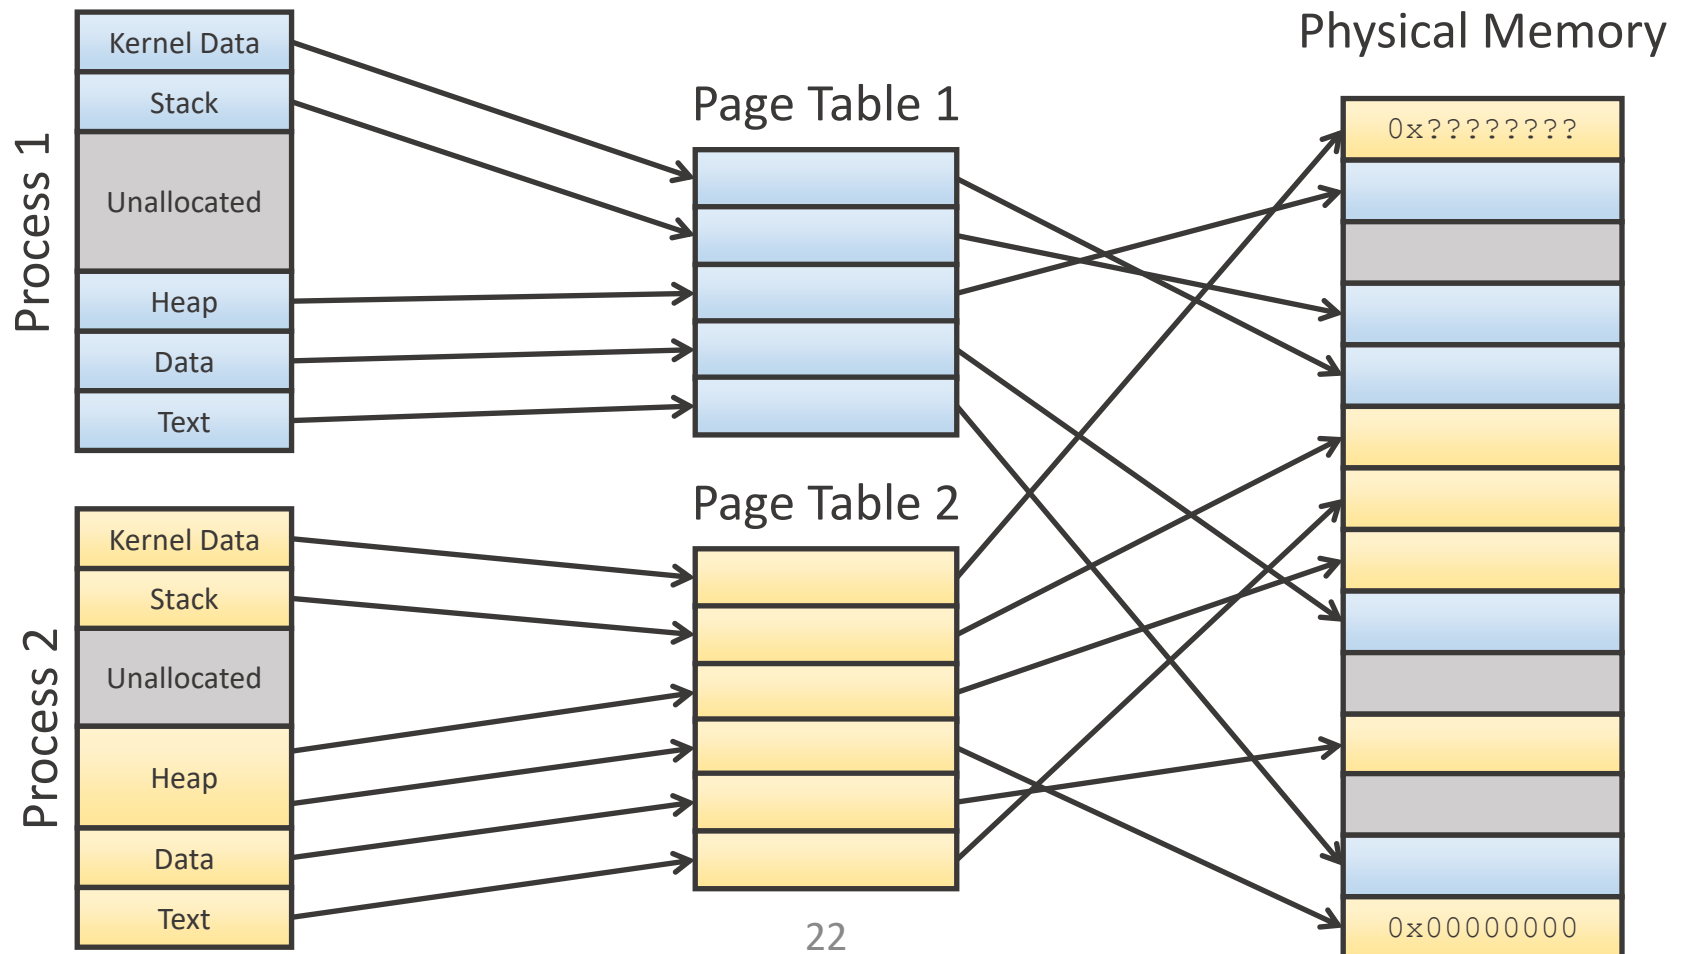
Page Tables

- All processes see an individual linear address space
- Page tables map from a linear address space to the physical address space



Page Tables

- Processes are separated by the page system



Paging

- Page Tables have a valid / invalid bit
 - Valid pages have page numbers allocated to the currently executing process
 - Invalid pages are either non-existent (not in the page table) or are in the page table but belong to other processes
- Memory access to an invalid page results in a segmentation / page fault or bus error
 - Trap causes context switch to kernel
 - Kernel sends SIGSEGV or SIGBUS to process
 - Usual behaviour is for process to end

Protecting Memory

- OS Integrity – preserved by separation of users and kernel space
- Separation of users
 - File management – logical memory object
 - Memory management – physical memory object

Summary

- Reference Monitors
- Operating System Integrity
- Privilege Elevation
- Memory Protection
- Page Tables

