

# COMP4131: Data Modelling and Analysis

## Lecture 7: Gaussian Process Regression and Classification

Daokun Zhang

University of Nottingham Ningbo China

*daokun.zhang@nottingham.edu.cn*

31 March, 2025

- 1 Gaussian Process
- 2 Gaussian Process Regression
- 3 Gaussian Process Classification

# Gaussian Process

# Linear Regression

Recall the probabilistic analysis of the standard linear regression model with Gaussian noise

$$f(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}), \quad y = f(\mathbf{x}) + \varepsilon,$$

where  $\varepsilon$  is the additive noise that follows an i.i.d Gaussian distribution with zero mean and variance  $\sigma_n^2$ :

$$\varepsilon \sim \mathcal{N}(0, \sigma_n^2).$$

Given the training data  $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  and  $\mathbf{y} = \{y_1, \dots, y_N\}$ , we can use maximum likelihood estimation (MLE) to find the solution of  $\mathbf{w}$  as

$$\mathbf{w}^* = (\boldsymbol{\Phi}^T \boldsymbol{\Phi})^{-1} \boldsymbol{\Phi}^T \mathbf{y},$$

where  $\boldsymbol{\Phi} = [\boldsymbol{\phi}(\mathbf{x}_1), \dots, \boldsymbol{\phi}(\mathbf{x}_N)]^T$  is an  $N \times M$  matrix.

# Bayesian Linear Regression

In the Bayesian formalism, we can specify a prior over the parameters, expressing our beliefs about the parameters before we look at the observations. We put a Gaussian prior with mean vector  $\mu_{\mathbf{w}}$  and covariance matrix  $\Sigma_{\mathbf{w}}$  on the weights

$$\mathbf{w} \sim \mathcal{N}(\mu_{\mathbf{w}}, \Sigma_{\mathbf{w}}).$$

Due to the randomness in  $\mathbf{w}$ ,  $f(x) = \mathbf{w}^T \phi(\mathbf{x})$  is no longer a deterministic but a **random function**. Let  $f_p$  denote the linear function value  $f(\mathbf{x}_n)$  at sample  $\mathbf{x}_n$ , then  $f_1, \dots, f_N$  follow a multivariate Gaussian distribution with

$$\begin{aligned}\mathbb{E}[f_p] &= \mu_{\mathbf{w}}^T \phi(\mathbf{x}_p), \\ \text{Cov}(f_n, f_m) &= \phi(\mathbf{x}_n)^T \Sigma_{\mathbf{w}} \phi(\mathbf{x}_m).\end{aligned}$$

Now we get a **Gaussian process**  $f(\mathbf{x})$  with mean function  $m(\mathbf{x}) = \mu_{\mathbf{w}}^T \phi(\mathbf{x})$  and covariance function  $k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \Sigma_{\mathbf{w}} \phi(\mathbf{x}')$ .

**Definition:** *A Gaussian process is a collection of random variables, any finite number of which have a joint Gaussian distribution.*

## Implications

- “a collection of random variables” means a set that has
  - a finite number of random variables
  - an infinite number of random variables
- “any finite number of which have a joint Gaussian distribution” means:
  - multiple random variables follow a joint multivariate Gaussian distribution
  - after marginalization, a single random variable follows a univariate Gaussian distribution

# Gaussian Process

A Gaussian process is completely specified by its **mean function** and **covariance function**. We define mean function  $m(\mathbf{x})$  and the covariance function  $k(\mathbf{x}, \mathbf{x}')$  as

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})],$$
$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E} [\{f(\mathbf{x}) - m(\mathbf{x})\} \{f(\mathbf{x}') - m(\mathbf{x}')\}],$$

and write the Gaussian process as

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')).$$

Usually, for notational simplicity we will take the mean function to be zero, i.e., set  $m(\mathbf{x}) = 0$ .

**Example:** For Bayesian linear regression model  $f(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x})$  with prior  $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \mathbf{\Sigma}_{\mathbf{w}})$ , we have the mean and covariance function

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})] = \phi(\mathbf{x})^T \mathbb{E}[\mathbf{w}] = 0,$$
$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[f(\mathbf{x})f(\mathbf{x}')] = \phi(\mathbf{x})^T \mathbb{E}[\mathbf{w}\mathbf{w}^T] \phi(\mathbf{x}') = \phi(\mathbf{x})^T \mathbf{\Sigma}_{\mathbf{w}} \phi(\mathbf{x}').$$

# Kernel Trick

The covariance function  $k(\mathbf{x}, \mathbf{x}')$  can be extended to any **kernel functions**.

For all  $\mathbf{x}$  and  $\mathbf{x}'$  in the input space  $\mathcal{X}$ , the function  $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is called a kernel function, if there exists a feature map  $\varphi: \mathcal{X} \rightarrow \mathcal{V}$  that satisfies

$$k(\mathbf{x}, \mathbf{x}') = \langle \varphi(\mathbf{x}), \varphi(\mathbf{x}') \rangle_{\mathcal{V}},$$

where  $\mathcal{V}$  is a inner product space, and  $\langle \cdot, \cdot \rangle_{\mathcal{V}}$  is the inner product operation defined by  $\mathcal{V}$ .

An alternative definition can be formulated by the **positive semidefinite (PSD)** property: For any points  $(\mathbf{x}_1, \dots, \mathbf{x}_n)$  in  $\mathcal{X}$ , and all choices of  $n$  real-valued coefficients  $(c_1, \dots, c_n)$ , the function  $k: \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is a kernel function, if

$$\sum_{i=1}^n \sum_{j=1}^n k(\mathbf{x}_i, \mathbf{x}_j) c_i c_j \geq 0.$$

The spanned matrix,  $\mathbf{K} \in \mathbb{R}^n \times \mathbb{R}^n$  with its  $ij$ -th entry  $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ , is called **Gram matrix**.



# Kernel Trick

Many kernels can be chosen for various application scenarios

- Fisher kernel
- Polynomial kernel
- Radial basis function kernel (RBF)
- String kernels
- Graph kernels

The Radial basis function kernel (RBF) is also called squared-exp or Gaussian kernel, which is formulated as

$$k_{\text{SE}}(\mathbf{x}_i, \mathbf{x}_j) = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|_2^2}{2\ell^2}\right).$$

# Gaussian Process Regression

# Prediction with Noise-free Observations

We first consider the ideal case where the observation  $y$  can be described by a Gaussian process  $f(\mathbf{x}) \sim \mathcal{GP}(0, k(\mathbf{x}, \mathbf{x}'))$  in a **noise-free** manner, i.e.,  $y = f(\mathbf{x})$ . For a set of test samples with size  $n_*$ , and feature matrix  $\mathbf{X}_*$ , its function value vector  $\mathbf{f}_*$  follows a multivariate Gaussian distribution

$$\mathbf{f}_* \sim \mathcal{N}(\mathbf{0}, \mathbf{K}(\mathbf{X}_*, \mathbf{X}_*)),$$

where  $\mathbf{K}(\cdot, \cdot)$  is the covariance matrix between two collections of input feature vectors.


If we are given a set of training samples with size  $n$ , and feature matrix  $\mathbf{X}$ , its function value vector  $\mathbf{f}$  should follow a multivariate Gaussian distribution jointly with  $\mathbf{f}_*$

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} \mathbf{K}(\mathbf{X}, \mathbf{X}) & \mathbf{K}(\mathbf{X}, \mathbf{X}_*) \\ \mathbf{K}(\mathbf{X}_*, \mathbf{X}) & \mathbf{K}(\mathbf{X}_*, \mathbf{X}_*) \end{bmatrix}\right).$$


# Prediction with Noise-free Observations

Given the observations of  $\mathbf{f}$ , the **posterior distribution** of  $\mathbf{f}_*$  is also a Gaussian distribution

$$\mathbf{f}_* | \mathbf{X}_*, \mathbf{X}, \mathbf{f} \sim \mathcal{N} \left( \mathbf{K}(\mathbf{X}_*, \mathbf{X}) \mathbf{K}(\mathbf{X}, \mathbf{X})^{-1} \mathbf{f} \right. \\ \left. \mathbf{K}(\mathbf{X}_*, \mathbf{X}_*) - \mathbf{K}(\mathbf{X}_*, \mathbf{X}) \mathbf{K}(\mathbf{X}, \mathbf{X})^{-1} \mathbf{K}(\mathbf{X}, \mathbf{X}_*) \right).$$



(a), prior



(b), posterior

The prior and posterior distributions of  $\mathbf{f}_*$  with the RBF kernel function  $k(\mathbf{x}_p, \mathbf{x}_q) = \exp(-\frac{1}{2}\|\mathbf{x}_p - \mathbf{x}_q\|_2^2)$ .

# Prediction using Noisy Observations

It is typical for more realistic modelling situations that we do not have access to function values themselves, but only **noisy versions** thereof  $y = f(\mathbf{x}) + \varepsilon$ .

Assuming additive **independent identically distributed** Gaussian noise  $\varepsilon$  with zero mean and variance  $\sigma_n^2$ , the prior on the noisy observations becomes another Gaussian distribution with mean  $\mathbb{E}[y] = \mathbb{E}[f(\mathbf{x}) + \varepsilon] = 0$ , and covariance

$$\text{cov}(y_p, y_q) = k(\mathbf{x}_p, \mathbf{x}_q) + \sigma_n^2 \delta_{pq} \text{ or } \text{cov}(\mathbf{y}) = K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I},$$

where  $\delta_{pq}$  is a Kronecker delta symbol which is one if and only if  $p = q$  and zero otherwise.

# Gaussian Process Regression

We can write the joint distribution of the observed target values  $\mathbf{y}$  and the function values  $\mathbf{f}_*$  of test samples under the prior as

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left( \mathbf{0}, \begin{bmatrix} \mathbf{K}(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I} & \mathbf{K}(\mathbf{X}, \mathbf{X}_*) \\ \mathbf{K}(\mathbf{X}_*, \mathbf{X}) & \mathbf{K}(\mathbf{X}_*, \mathbf{X}_*) \end{bmatrix} \right).$$


Deriving the conditional distribution, we arrive at the [key predictive equations](#) for [Gaussian process regression](#)

$$\mathbf{f}_* | \mathbf{X}, \mathbf{y}, \mathbf{X}_* \sim \mathcal{N}(\bar{\mathbf{f}}_*, \text{cov}(\mathbf{f}_*)), \text{ where}$$

$$\bar{\mathbf{f}}_* \triangleq \mathbb{E}[\mathbf{f}_* | \mathbf{X}, \mathbf{y}, \mathbf{X}_*] = \mathbf{K}(\mathbf{X}_*, \mathbf{X}) [\mathbf{K}(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I}]^{-1} \mathbf{y},$$

$$\text{cov}(\mathbf{f}_*) = \mathbf{K}(\mathbf{X}_*, \mathbf{X}_*) - \mathbf{K}(\mathbf{X}_*, \mathbf{X}) [\mathbf{K}(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I}]^{-1} \mathbf{K}(\mathbf{X}, \mathbf{X}_*).$$

# Gaussian Process Regression



Graphical model (chain graph) for a Gaussian process for regression. Squares represent **observed variables** and circles represent **unknowns**.

# Gaussian Process Regression

By using a compact form of the notation setting  $\mathbf{K} = \mathbf{K}(\mathbf{X}, \mathbf{X})$  and  $\mathbf{k}_* = \mathbf{K}(\mathbf{X}, \mathbf{x}_*)$  for a test sample  $\mathbf{x}_*$ , the prediction  $f_*$  for the single test sample  $\mathbf{x}_*$  has mean  $\bar{f}_*$  and variance  $\mathbb{V}[f_*]$  as

$$\bar{f}_* = \mathbf{k}_*^\top (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y},$$
$$\mathbb{V}[f_*] = k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{k}_*^\top (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k}_*.$$

Another way to look at the expression of the mean  $\bar{f}_*$  is to see it as a **linear combination** of  $N$  kernel functions, each one centered on a training point, by writing

$$\bar{f}_* = \sum_{i=1}^N \alpha_i k(\mathbf{x}_i, \mathbf{x}_*),$$

where  $\boldsymbol{\alpha} = (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$ .



# Gaussian Process Regression

**input:**  $X$  (inputs),  $\mathbf{y}$  (targets),  $k$  (covariance function),  $\sigma_n^2$  (noise level),  $\mathbf{x}_*$  (test input)

2:  $L := \text{cholesky}(K + \sigma_n^2 I)$   
     $\boldsymbol{\alpha} := L^\top \backslash (L \backslash \mathbf{y})$

4:  $\bar{f}_* := \mathbf{k}_*^\top \boldsymbol{\alpha}$  } predictive mean  
     $\mathbf{v} := L \backslash \mathbf{k}_*$  }

6:  $\mathbb{V}[f_*] := k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{v}^\top \mathbf{v}$  } predictive variance  
     $\log p(\mathbf{y}|X) := -\frac{1}{2} \mathbf{y}^\top \boldsymbol{\alpha} - \sum_i \log L_{ii} - \frac{n}{2} \log 2\pi$

8: **return:**  $\bar{f}_*$  (mean),  $\mathbb{V}[f_*]$  (variance),  $\log p(\mathbf{y}|X)$  (log marginal likelihood)

Algorithm for Gaussian process regression, where the predictions (the target value's mean and variance) at the test sample  $\mathbf{x}_*$  and the [log marginal likelihood](#) on training set are returned.

# Varying the Hyperparameters

Setting the RBF kernel  $k(\mathbf{x}_p, \mathbf{x}_q) = \sigma_f^2 \exp(-\frac{1}{2\ell^2} \|\mathbf{x}_p - \mathbf{x}_q\|_2^2)$  as the covariance function, the covariance between the observations  $y_p$  and  $y_q$  becomes


$$\text{cov}(y_p, y_q) = \sigma_f^2 \exp(-\frac{1}{2\ell^2} \|\mathbf{x}_p - \mathbf{x}_q\|_2^2) + \sigma_n^2 \delta_{pq},$$

which implies that the model relies on three hyperparameters:


- signal variance  $\sigma_f^2$
- noise variance  $\sigma_n^2$
- length-scale  $\ell$

The log marginal likelihood can be used to tune the hyperparameters.


# Varying the Hyperparameters



(a),  $\ell = 1$



(b),  $\ell = 0.3$



(c),  $\ell = 3$

Small length-scale  $\ell = 0.3 \rightarrow$  a sharply varying function  $f$  with small noise.

Large length-scale  $\ell = 3 \rightarrow$  a slowly varying function  $f$  with large noise.

Moderate length-scale  $\ell = 1 \rightarrow$  exact fitting of the underlying function  $f$ .

# Gaussian Process Classification

# Binary Gaussian Process Classifier

For binary classification with label candidate set  $\{+1, -1\}$ , linear regression uses the **sigmoid** function to squash the linear regression output into the label probability  $p(y = +1|\mathbf{x}) = \sigma(\mathbf{w}^T \phi(\mathbf{x}))$ .

In Gaussian process, the latent function value  $f(\mathbf{x})$  is used to construct the linear regression output.

Similarly, for **Gaussian process classification**, we can also construct the **label probability** via the **sigmoid** function

$$p(y = +1|f(\mathbf{x})) = \sigma(f(\mathbf{x})).$$

Different from  $\mathbf{w}^T \phi(\mathbf{x})$ ,  $f(\mathbf{x})$  is Gaussian random variable, we need to marginalize it to get the posterior probability

$$p(y = +1|\mathbf{x}) = \int p(y = +1|f)p(f|\mathbf{x})df.$$

# Binary Gaussian Process Classifier

Given a training set  $\{\mathbf{X}, \mathbf{y}\}$ , for a new test sample  $\mathbf{x}_*$ , Gaussian process classifier predicts its label by estimating the posterior probability

$$\begin{aligned}\bar{\pi}_* &\triangleq p(y_* = +1 | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) = \int p(y_* = +1 | f_*) p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) df_* \\ &= \int \sigma(f_*) p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) df_*.\end{aligned}$$

To achieve the estimation, we need first calculate the posterior distribution of the latent variable  $f_*$  corresponding to the test sample  $\mathbf{x}_*$

$$p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) = \int p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) p(\mathbf{f} | \mathbf{X}, \mathbf{y}) d\mathbf{f},$$

where  $p(\mathbf{f} | \mathbf{X}, \mathbf{y}) = p(\mathbf{y} | \mathbf{f}) p(\mathbf{f} | \mathbf{X}) / p(\mathbf{y} | \mathbf{X})$  is the posterior over the latent variables and  $p(\mathbf{y} | \mathbf{f})$  is computed by applying the elementwise sigmoid function to  $\mathbf{f}$ .

# Laplace Approximation for Binary GP Classifier

In the regression case,  $\mathbf{y}$  and  $f_*$  jointly follow a multivariate Gaussian distribution, the conditional probability distribution  $p(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*)$  can be computed analytically.

However, in the classification case, due to the computational intractability of the distribution  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$ , the posterior distribution  $p(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*)$  cannot be computed analytically with the intractable integral.

To solve this problem, we resort to the [Laplace approximation](#) that approximates the distribution  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  with a Gaussian distribution  $q(\mathbf{f}|\mathbf{X}, \mathbf{y})$ , obtained by doing a second order Taylor expansion of  $\log p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  around the maximum of the posterior

$$q(\mathbf{f}|\mathbf{X}, \mathbf{y}) = \mathcal{N}(\mathbf{f}|\hat{\mathbf{f}}, \mathbf{A}^{-1}) \propto \exp\left(-\frac{1}{2}(\mathbf{f} - \hat{\mathbf{f}})^\top \mathbf{A}(\mathbf{f} - \hat{\mathbf{f}})\right),$$

where  $\hat{\mathbf{f}} = \operatorname{argmax}_{\mathbf{f}} p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  and  $\mathbf{A} = -\nabla\nabla \log p(\mathbf{f}|\mathbf{X}, \mathbf{y})|_{\mathbf{f}=\hat{\mathbf{f}}}$  is the Hessian matrix of the negative log posterior at that point.

# Laplace Approximation for Binary GP Classifier

By Bayes' rule, the posterior over the latent variables  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  is given by

$$p(\mathbf{f}|\mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y}, \mathbf{f}|\mathbf{X})}{p(\mathbf{y}|\mathbf{X})} = \frac{p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\mathbf{X})}{p(\mathbf{y}|\mathbf{X})},$$

but  $p(\mathbf{y}|\mathbf{X})$  is independent of  $\mathbf{f}$ , we only need to consider the un-normalized posterior  $p(\mathbf{y}, \mathbf{f}|\mathbf{X}) = p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\mathbf{X})$  when maximizing w.r.t.  $\mathbf{f}$ .

Taking the logarithm on the un-normalized posterior, we get

$$\begin{aligned}\Psi(\mathbf{f}) &\triangleq \log p(\mathbf{y}|\mathbf{f}) + \log p(\mathbf{f}|\mathbf{X}) \\ &= \log p(\mathbf{y}|\mathbf{f}) - \frac{1}{2}\mathbf{f}^T \mathbf{K}^{-1}\mathbf{f} - \frac{1}{2}\log |\mathbf{K}| - \frac{n}{2}\log 2\pi,\end{aligned}$$

where  $p(\mathbf{f}|\mathbf{X})$  is a Gaussian distribution  $\mathcal{N}(\mathbf{0}, \mathbf{K})$ ,  $\mathbf{K}$  is short for the kernel matrix  $\mathbf{K}(\mathbf{X}, \mathbf{X})$ , and  $|\mathbf{K}|$  is the determinant of the kernel matrix  $\mathbf{K}$ .



# Laplace Approximation for Binary GP Classifier

Differentiating  $\Psi(\mathbf{f})$  w.r.t.  $\mathbf{f}$ , we obtain

$$\begin{aligned}\nabla \Psi(\mathbf{f}) &= \nabla \log p(\mathbf{y}|\mathbf{f}) - \mathbf{K}^{-1}\mathbf{f}, \\ \nabla \nabla \Psi(\mathbf{f}) &= \nabla \nabla \log p(\mathbf{y}|\mathbf{f}) - \mathbf{K}^{-1},\end{aligned}$$

where  $\nabla \nabla \log p(\mathbf{y}|\mathbf{f})$  is diagonal, since the distribution for  $y_i = +1$  or  $-1$  depends only on  $f_i$ , not on  $f_{j \neq i}$ .

We can apply the [Newton-Raphson](#) method to find the latent function value vector  $\hat{\mathbf{f}}$  that maximizes  $\Psi(\mathbf{f})$ , and obtain the approximated probability distribution as

$$q(\mathbf{f}|\mathbf{X}, \mathbf{y}) = \mathcal{N}\left(\hat{\mathbf{f}}, (\mathbf{K}^{-1} + \mathbf{W})^{-1}\right),$$

where  $\mathbf{W} \triangleq -\nabla \nabla \log p(\mathbf{y}|\mathbf{f})|_{\mathbf{f}=\hat{\mathbf{f}}}$ .

# Laplace Approximation for Binary GP Classifier

With  $q(\mathbf{f}|\mathbf{X}, \mathbf{y})$  approximating  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$ , we can get the approximation for the posterior  $p(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*)$  as

$$q(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*) = \int p(f_*|\mathbf{X}, \mathbf{x}_*, \mathbf{f}) q(\mathbf{f}|\mathbf{X}, \mathbf{y}) d\mathbf{f},$$

where both  $p(f_*|\mathbf{X}, \mathbf{x}_*, \mathbf{f})$  and  $q(\mathbf{f}|\mathbf{X}, \mathbf{y})$  are Gaussian distributions. It can be proved that  $q(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*)$  is also a Gaussian distribution with mean and variance as

$$\mathbb{E}_q[f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*] = \mathbf{k}_*^T \mathbf{K}^{-1} \hat{\mathbf{f}},$$

$$\mathbb{V}_q[f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*] = k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{k}_*^T (\mathbf{K} + \mathbf{W}^{-1})^{-1} \mathbf{k}_*.$$

Finally, the label probability for the test sample  $\mathbf{x}_*$  can be computed as

$$\bar{\pi}_* \triangleq p(y_* = +1|\mathbf{X}, \mathbf{y}, \mathbf{x}_*) := \int \sigma(f_*) q(f_*|\mathbf{X}, \mathbf{y}, \mathbf{x}_*) df_*.$$

Numeric integration techniques are generally used to calculate the one-dimensional integral.

# Binary Gaussian Process Classifier

```
input:  $K$  (covariance matrix),  $\mathbf{y}$  ( $\pm 1$  targets),  $p(\mathbf{y}|\mathbf{f})$  (likelihood function)
2:  $\mathbf{f} := \mathbf{0}$  initialization
   repeat Newton iteration
4:    $W := -\nabla \nabla \log p(\mathbf{y}|\mathbf{f})$ 
      $L := \text{cholesky}(I + W^{\frac{1}{2}} K W^{\frac{1}{2}})$ 
6:    $\mathbf{b} := W\mathbf{f} + \nabla \log p(\mathbf{y}|\mathbf{f})$ 
      $\mathbf{a} := \mathbf{b} - W^{\frac{1}{2}} L^{\top} \setminus (L \setminus (W^{\frac{1}{2}} K \mathbf{b}))$ 
8:    $\mathbf{f} := K \mathbf{a}$ 
   until convergence
10:  $\log q(\mathbf{y}|\hat{\mathbf{X}}, \theta) := -\frac{1}{2} \mathbf{a}^{\top} \mathbf{f} + \log p(\mathbf{y}|\mathbf{f}) - \sum_i \log L_{ii}$ 
   return:  $\hat{\mathbf{f}} := \mathbf{f}$  (post. mode),  $\log q(\mathbf{y}|\hat{\mathbf{X}}, \theta)$  (approx. log marg. likelihood)
```

Training algorithm for the binary Gaussian process classifier with Laplace approximation. The mode  $\hat{\mathbf{f}}$  for the posterior probability distribution  $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$  is returned, as well as the **approximated log marginal likelihood**, which can be use to tune hyperparameters.

# Binary Gaussian Process Classifier

**input:**  $\hat{\mathbf{f}}$  (mode),  $X$  (inputs),  $\mathbf{y}$  ( $\pm 1$  targets),  $k$  (covariance function),  
 $p(\mathbf{y}|\mathbf{f})$  (likelihood function),  $\mathbf{x}_*$  test input

2:  $W := -\nabla \nabla \log p(\mathbf{y}|\hat{\mathbf{f}})$

$L := \text{cholesky}(I + W^{\frac{1}{2}} K W^{\frac{1}{2}})$

4:  $\bar{f}_* := \mathbf{k}(\mathbf{x}_*)^\top \nabla \log p(\mathbf{y}|\hat{\mathbf{f}})$

$\mathbf{v} := L \backslash (W^{\frac{1}{2}} \mathbf{k}(\mathbf{x}_*))$

6:  $\mathbb{V}[f_*] := k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{v}^\top \mathbf{v}$

$\bar{\pi}_* := \int \sigma(z) \mathcal{N}(z|\bar{f}_*, \mathbb{V}[f_*]) dz$

8: **return:**  $\bar{\pi}_*$  (predictive class probability (for class 1))

Prediction algorithm for the binary Gaussian classifier with Laplace approximation.

# Multi-Class Gaussian Process Classifier

For the multi-class classification problem, where the class label  $y$  belongs to the label candidate set  $\{1, \dots, C\}$  with size  $C$ , we need to construct a Gaussian process for each class  $c \in \{1, \dots, C\}$ , with the latent function value  $f^c(\mathbf{x})$  evaluated at  $\mathbf{x}$ .

The conditional label probability  $p(y = c | f^1(\mathbf{x}), \dots, f^C(\mathbf{x}))$  can be constructed by the softmax function

$$p(y = c | f^1(\mathbf{x}), \dots, f^C(\mathbf{x})) = \frac{\exp(f^c(\mathbf{x}))}{\sum_{c'=1}^C \exp(f^{c'}(\mathbf{x}))}.$$

Following the same procedure as the binary classification case, given a test sample  $\mathbf{x}_*$ , we can estimate the posterior probabilities  $p(y_* = c | \mathbf{X}, \mathbf{y}, \mathbf{x}_*)$  to predict  $\mathbf{x}_*$ 's class label.

Williams CK, Rasmussen CE. **Gaussian Processes for Machine Learning**. Cambridge, MA: MIT press; 2006.

- Section 1 - Section 3

# The End